

# Fundamentals of alignment-free sequence analysis: *k*-mer hashing



Part II: Hashing strategies,  
hardware considerations, summary

Jens Zentgraf & Sven Rahmann  
ACM-BCB 2020



# Hashing

- Set ("universe")  $U$  of possible keys
- Set of keys  $K \subseteq U$  to be stored,  $|K| = N$
- Hash table (array) with  $P$  slots
- Hash function  $h$ 
  - $h: U \rightarrow \{0, \dots, P - 1\}$

# Hashing

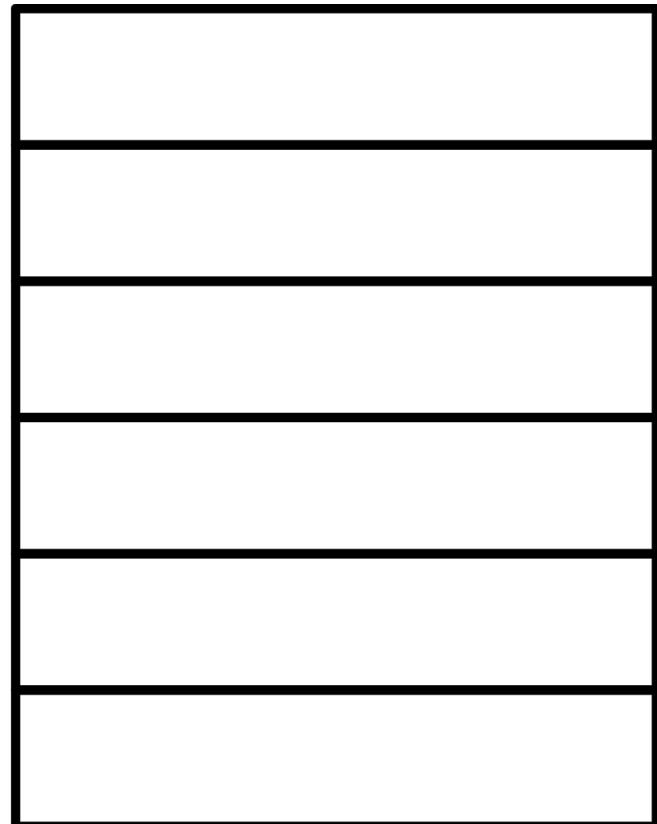
- Universe  $U$ : all possible first names (of finite length)
- Set of keys  $K = \{Anna, Franz, Bea, Fritz\}$
- Hash table with 6 slots
- Some function  $h$  mapping  $U$  to  $\{0, 1, 2, 3, 4, 5\}$ .

# Hashing

$h("Anna") = 3$

$h("Franz") = 5$

$h("Bea") = 2$

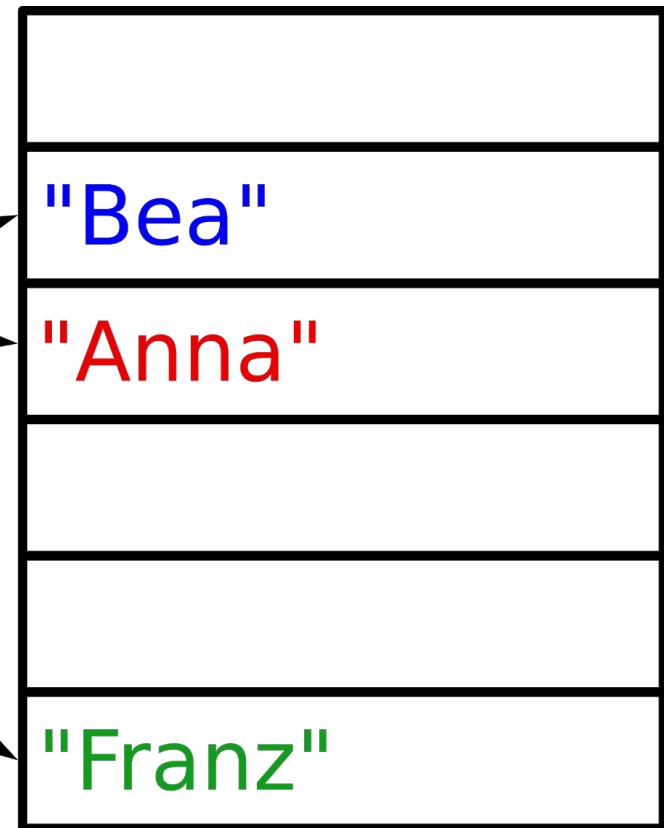


# Hashing

$$h("Anna") = 3$$

$$h("Franz") = 5$$

$$h("Bea") = 2$$



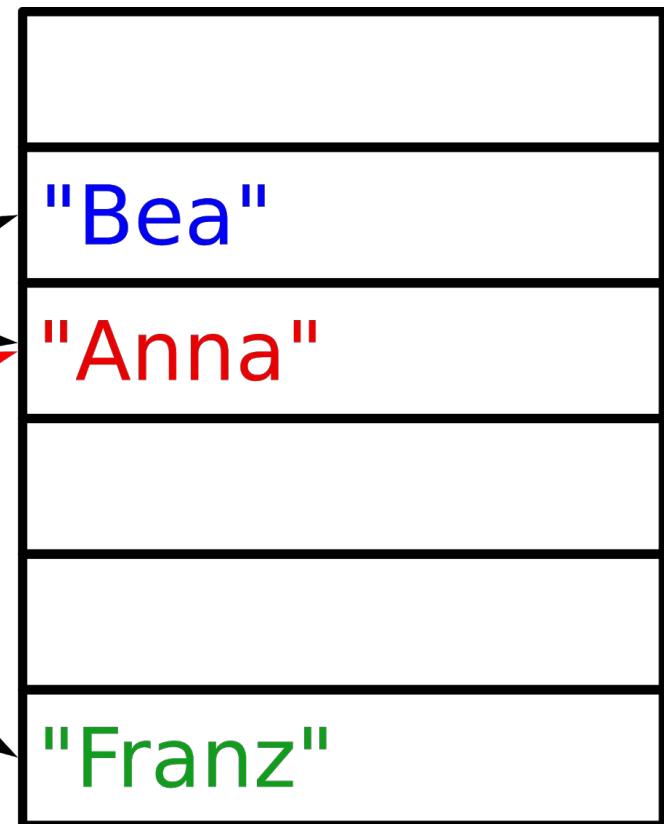
# Hashing

$$h("Anna") = 3$$

$$h("Franz") = 5$$

$$h("Bea") = 2$$

$$h("Fritz") = 3$$



# Collisions and strategies of collision resolution

Definition: collision

- Two elements are hashed to the same location
- $h(a) = h(b)$  for some  $a \neq b$

Strategies of collision resolution:

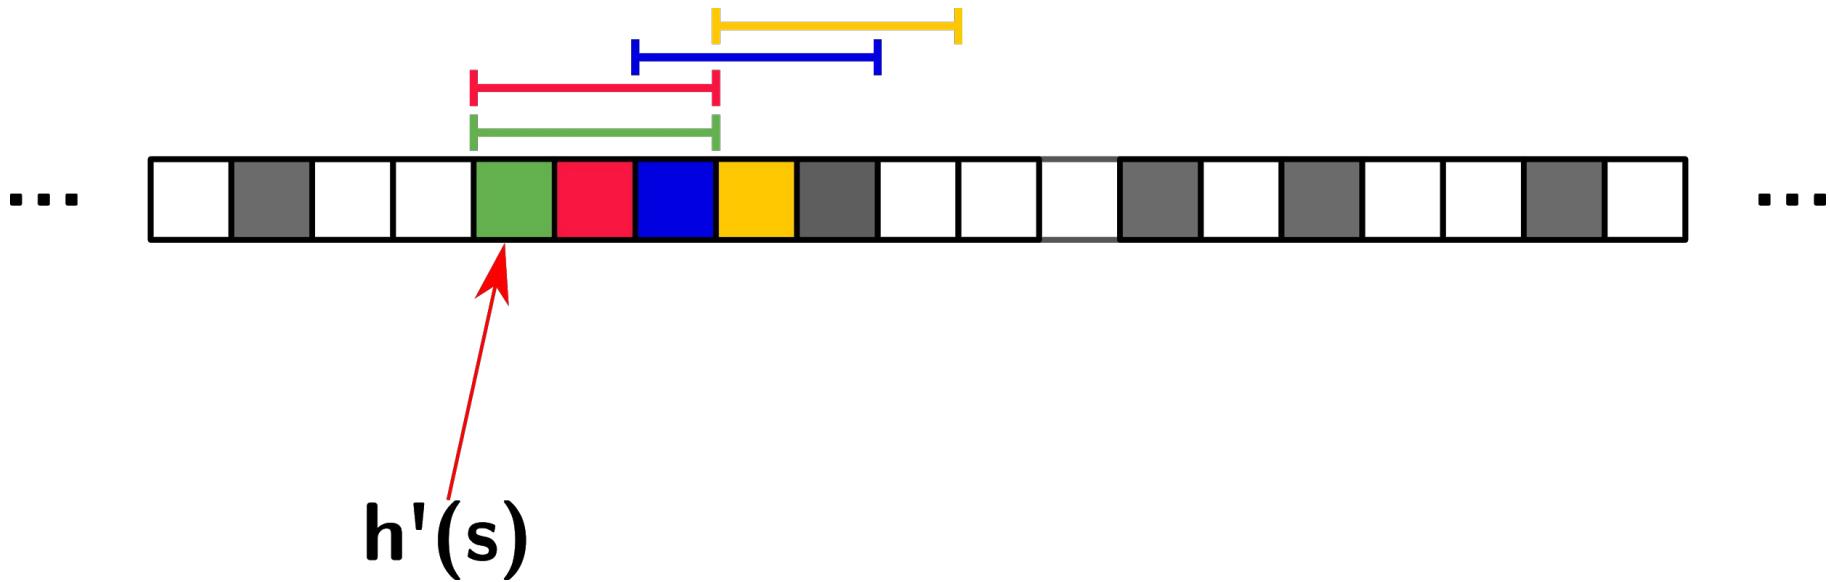
- Chaining (also: separate chaining)
- Open addressing (also: closed hashing)

# Hopscotch hashing

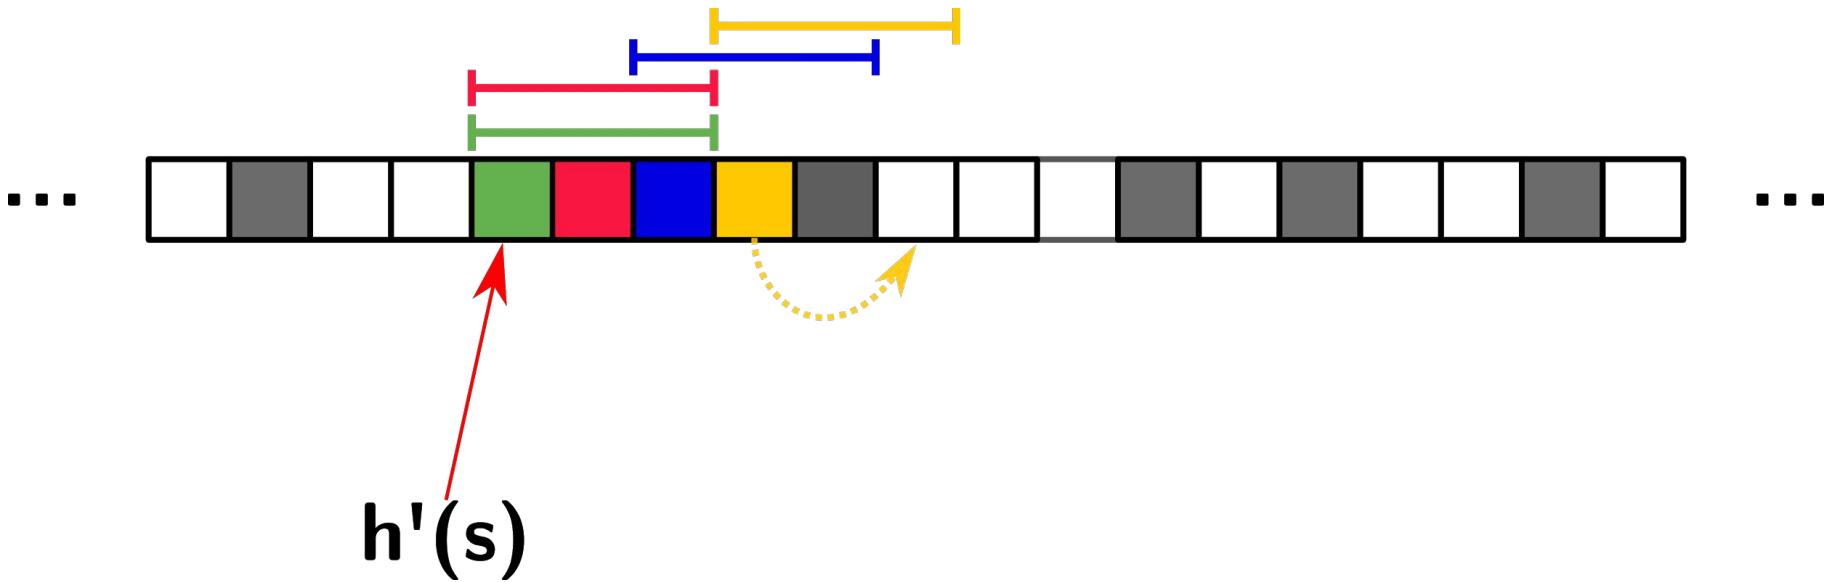
- One hash function  $h(x)$
- Maximum range  $r$  to insert element
- Element can be inserted between  $h(x)$  and  $h(x)+r-1$
- Bit array of length  $r$  for each position:  
Which elements in  $[h(x), h(x)+r-1]$  belong to  $h(x)$  ?
- Insert:  $O(N)$
- Lookup:  $O(r)$
- Compute hash position  $h(x)$
- If  $h(x)$  is occupied:
  - Find next empty slot  $y \geq h(x)$   
(can be far away)
  - Find an element that can be moved to the empty slot:
    - Start at position  $i = y - r - 1$
    - Check if element can be moved to  $y$
    - If not, repeat for increasing  $i$  up to  $y-1$
    - Fail if no element found.
  - Repeat with  $y = i$  (new empty slot), till  $x$  can be inserted in  $[h(x), h(x)+r-1]$ .

# Hopscotch hashing

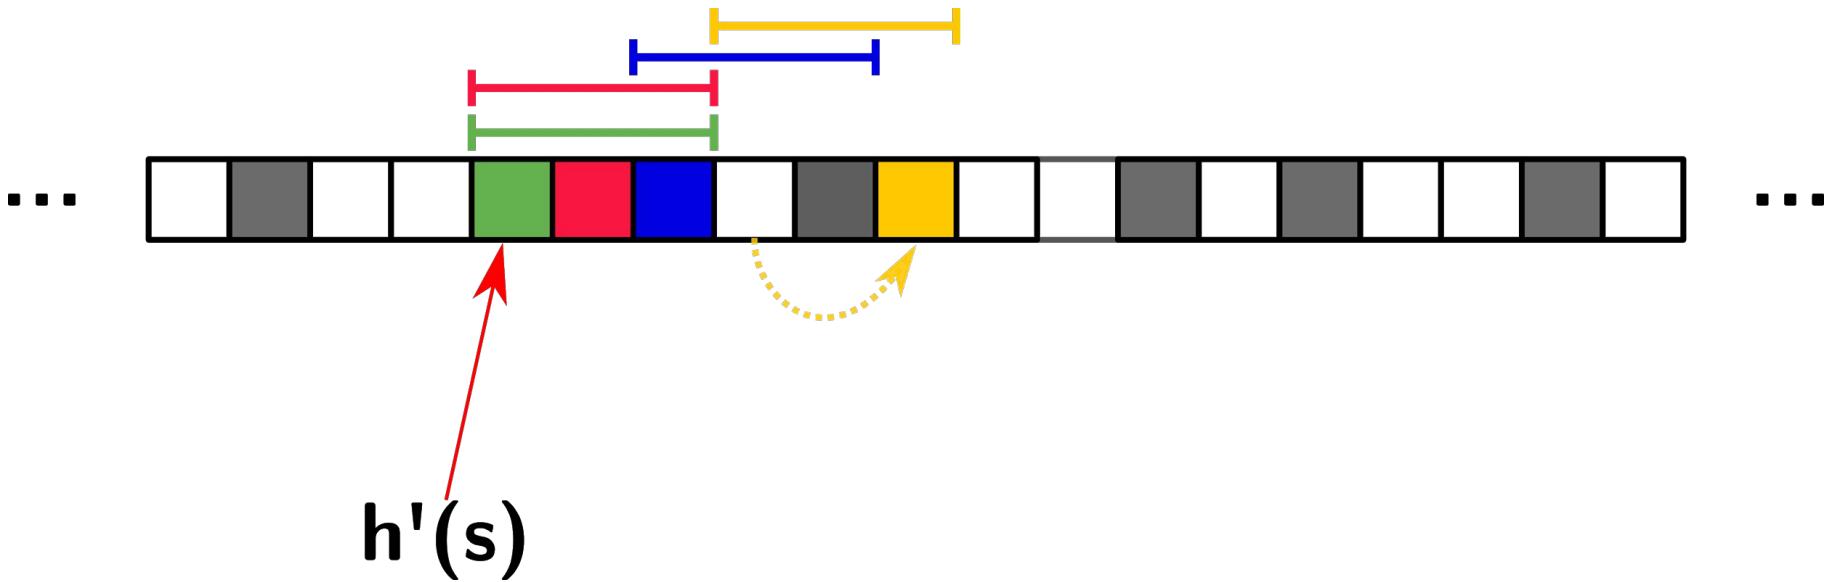
Example:  $r = 3$



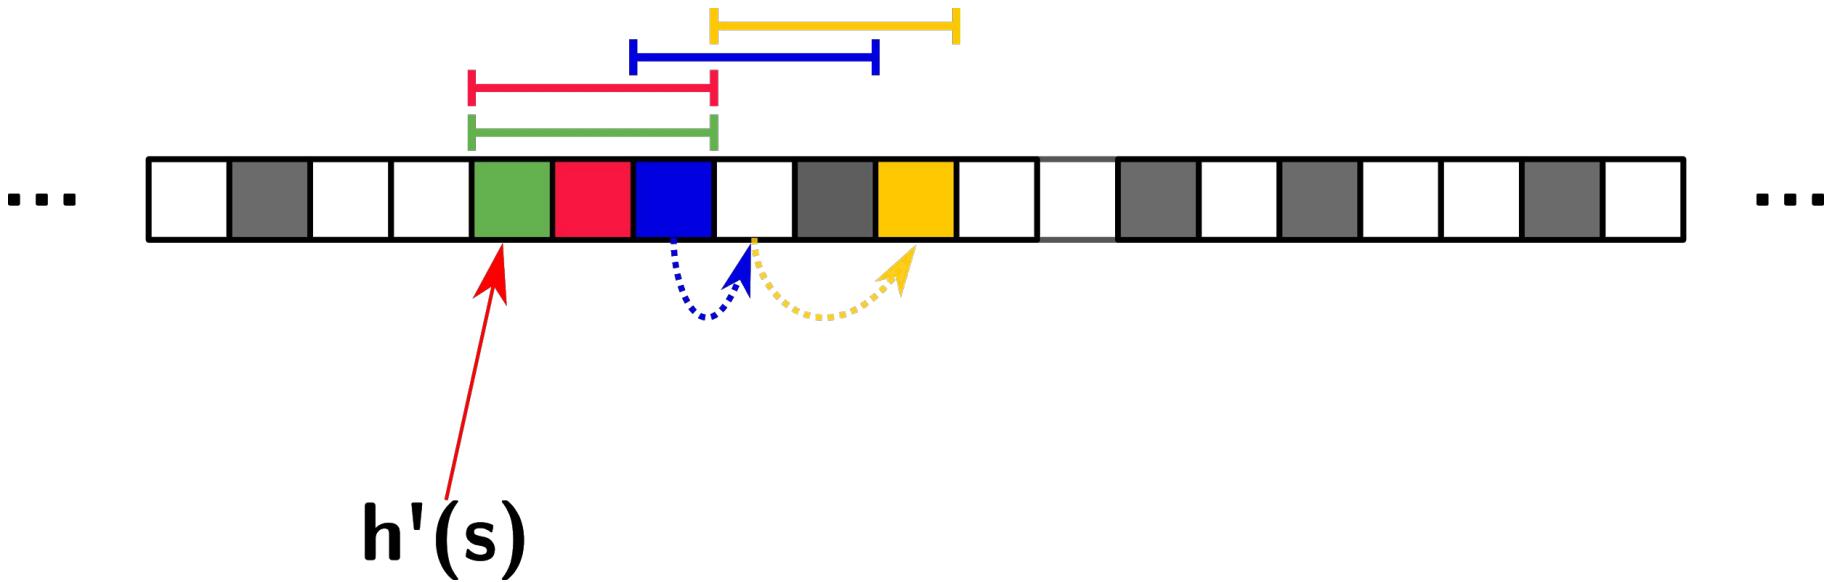
# Hopscotch hashing



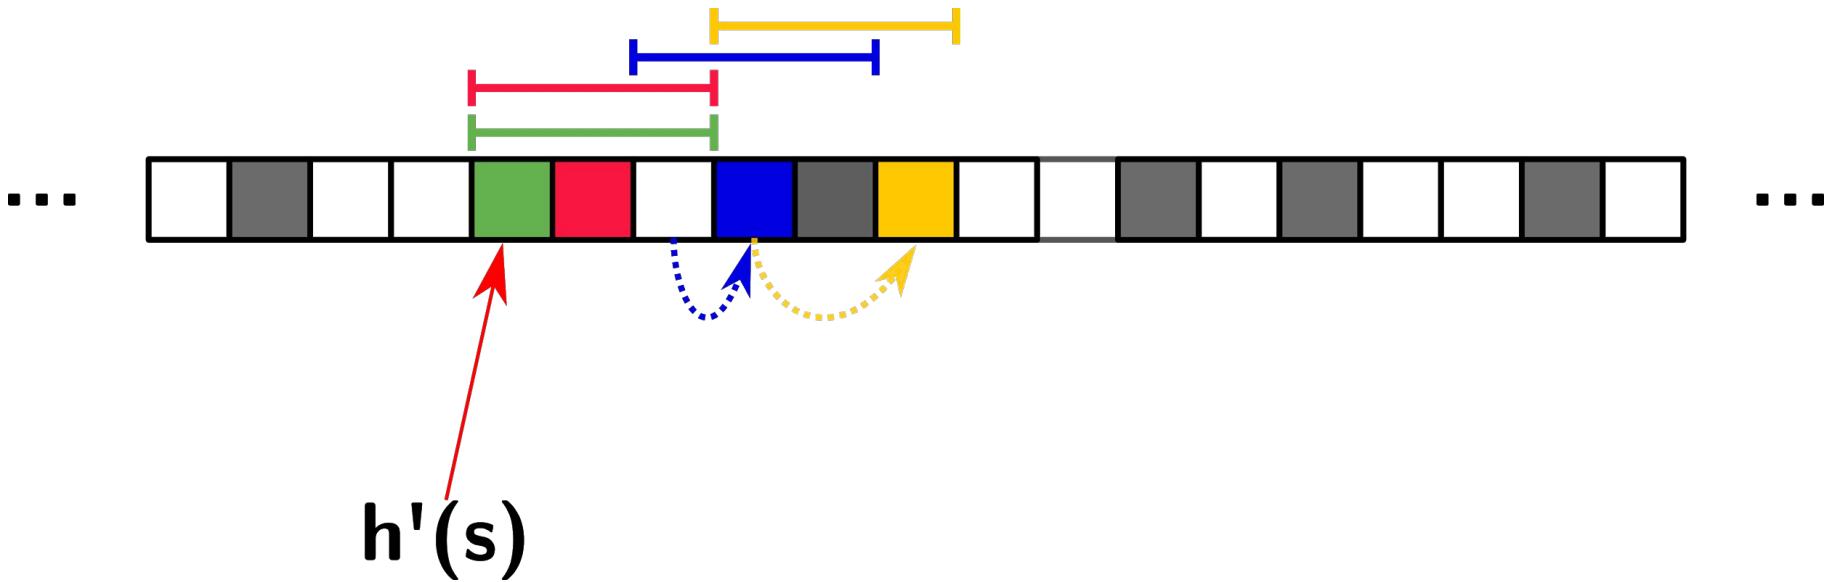
# Hopscotch hashing



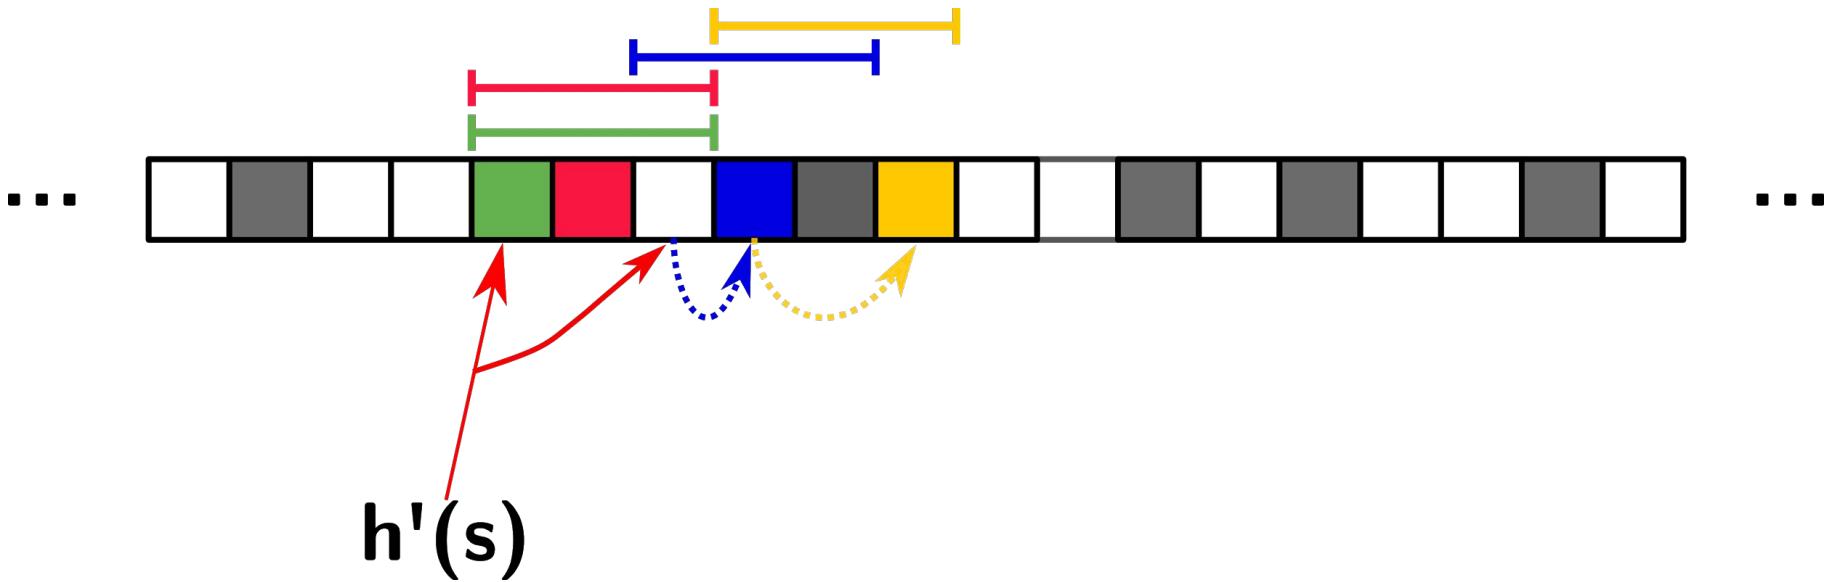
# Hopscotch hashing



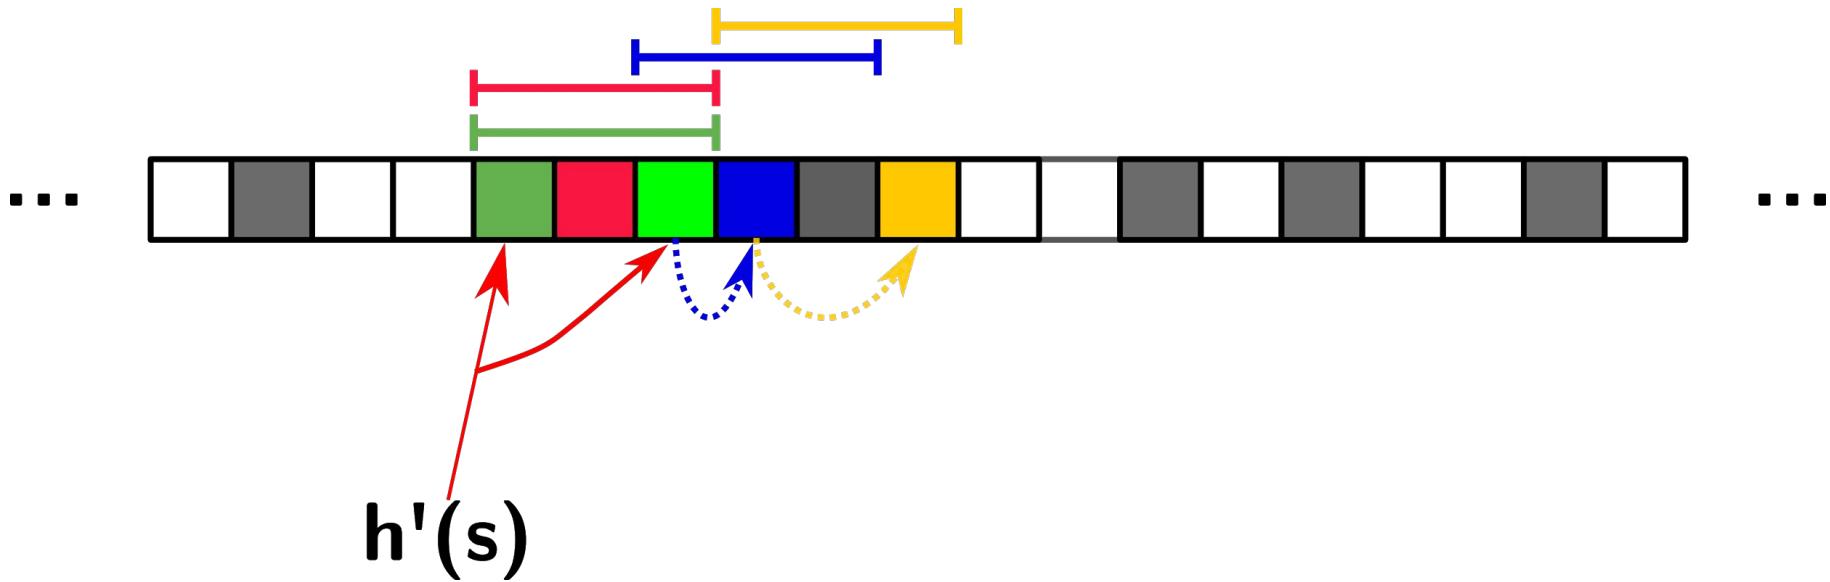
# Hopscotch hashing



# Hopscotch hashing



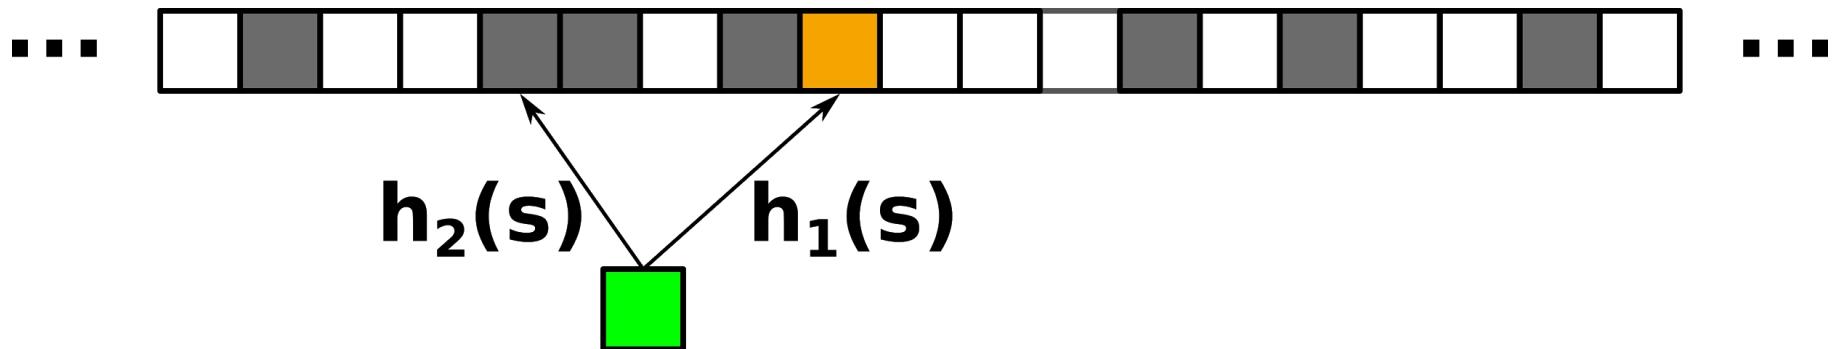
# Hopscotch hashing



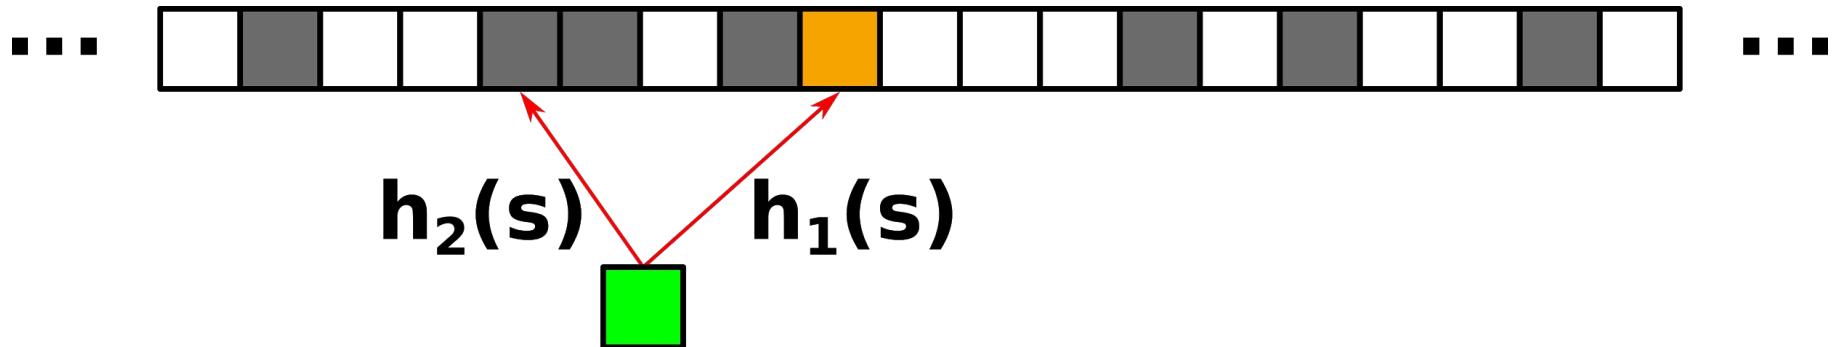
# Cuckoo hashing

- Two hash functions  $h_1$  and  $h_2$
- Insert new element  $x$  with  $h_1$ 
  - If position is occupied with element  $y$ , displace element  $y$
  - Insert  $y$  with alternative hash function
- Parameter  $w$ : Maximum number of displacements
  
- Insert:  $O(w)$
- Lookup:  $2 \in O(1)$

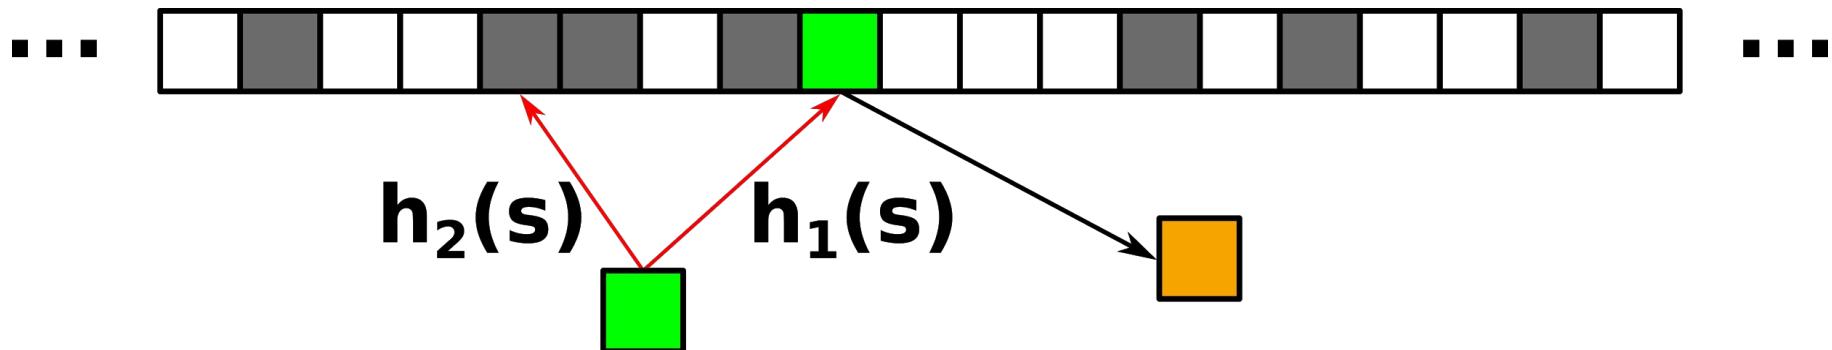
# Cuckoo hashing



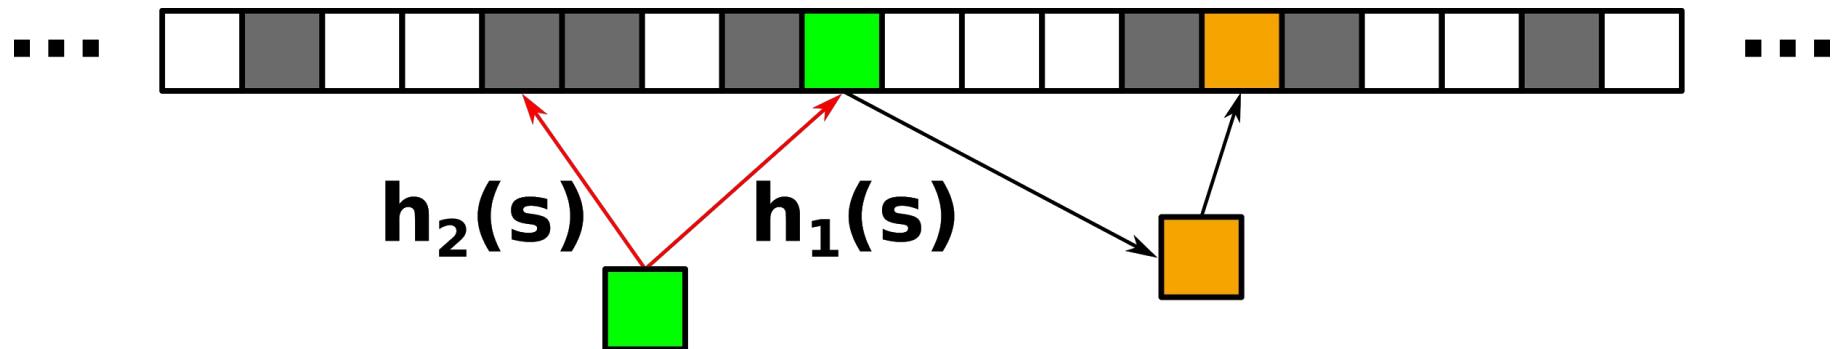
# Cuckoo hashing



# Cuckoo hashing



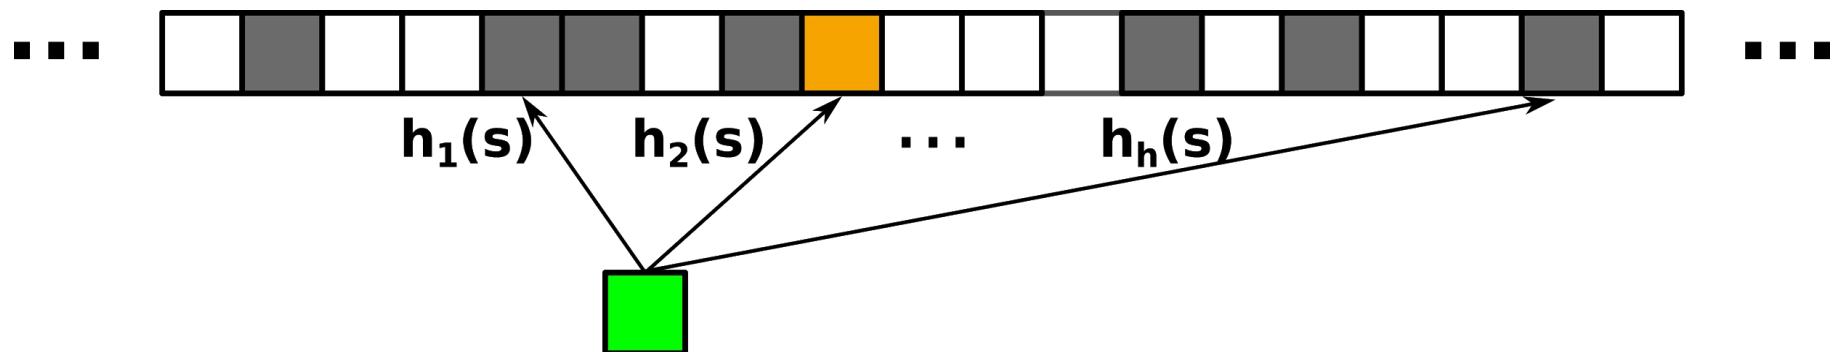
# Cuckoo hashing



# Cuckoo hashing with $h$ hash functions

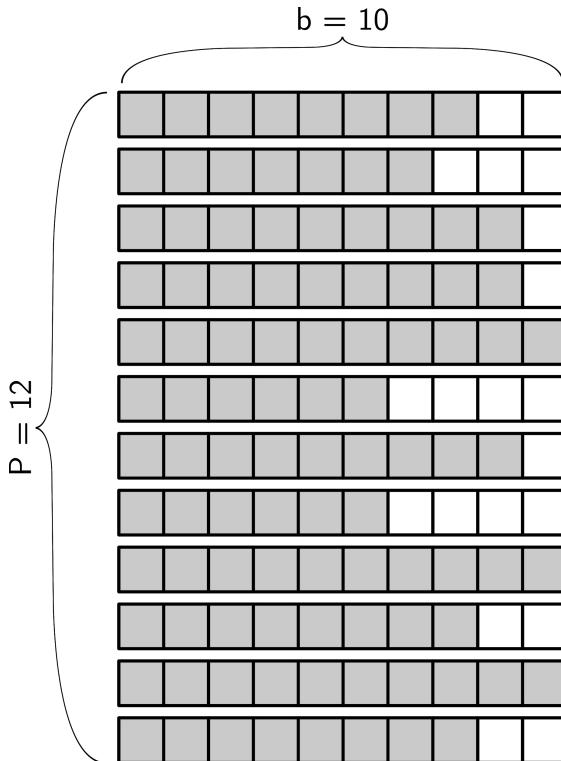
- Use  $h \geq 2$  hash functions
  - Better (more uniform) distribution of the elements
- Different displacement strategies:
  - random walk
  - LSA

# Cuckoo hashing with $h$ hash functions



# Cuckoo hashing with buckets

- Each position can store up to  $b$  elements ("buckets", "pages", "bins" of size  $b$ )
- Example:  $P = 12$  buckets of size  $b = 10$ : hash table with 120 slots.
- Must search bucket to find element  $x$ , until found or empty slot encountered, or entire bucket has been searched



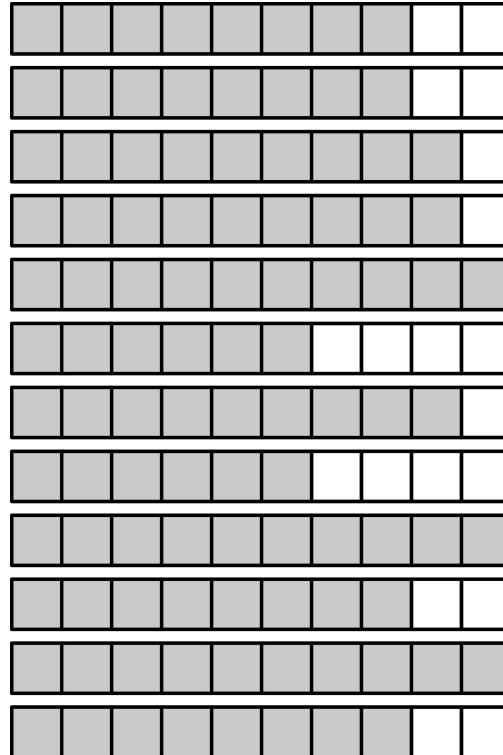
# (h,b) Cuckoo hashing

- Combination of
  - buckets with size b and
  - h hash functions
- Different insertion strategies
  - Random walk
  - Breadth first search
  - LSA<sub>max</sub>
  - Cuckopt
  - ...

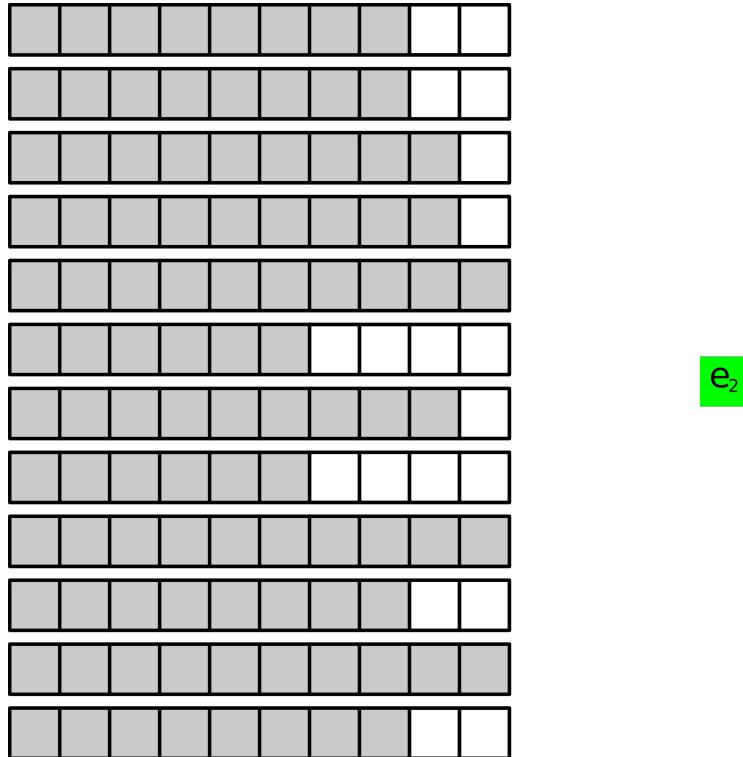
## (h,b) Cuckoo hashing: Random walk

- Check all hash functions for an empty slot
- If no empty slot exists
  - Choose one element and replace it
  - Insert the replaced element
- Maximum of  $w$  replacements
  
- Insert:  $O(w)$
- Lookup:  $O(h)$

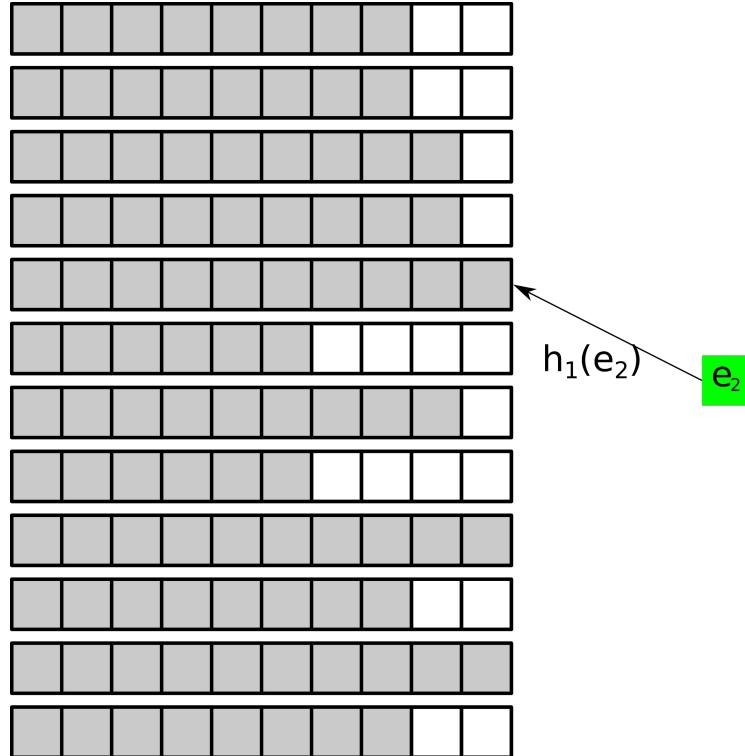
## (h,b) Cuckoo hashing: Random walk



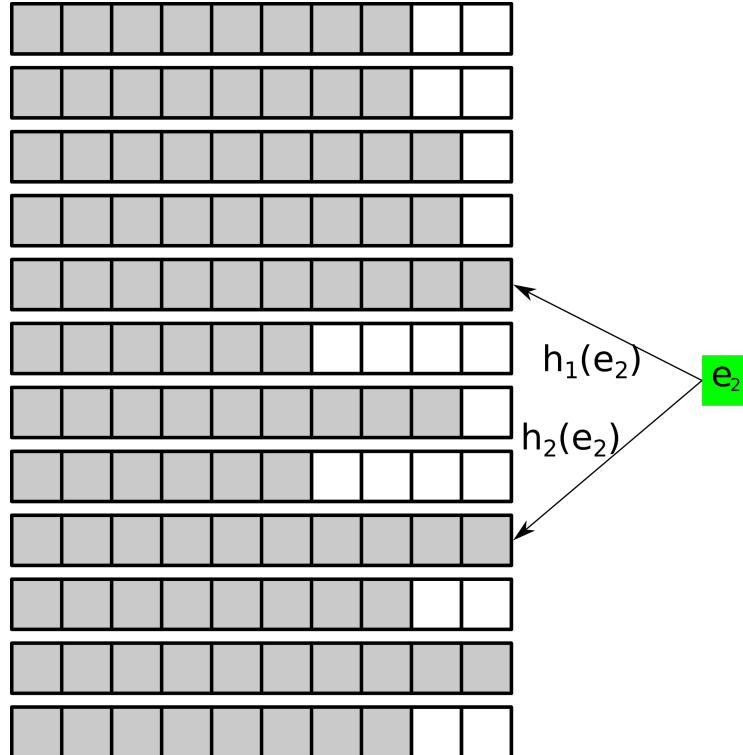
# (h,b) Cuckoo hashing: Random walk



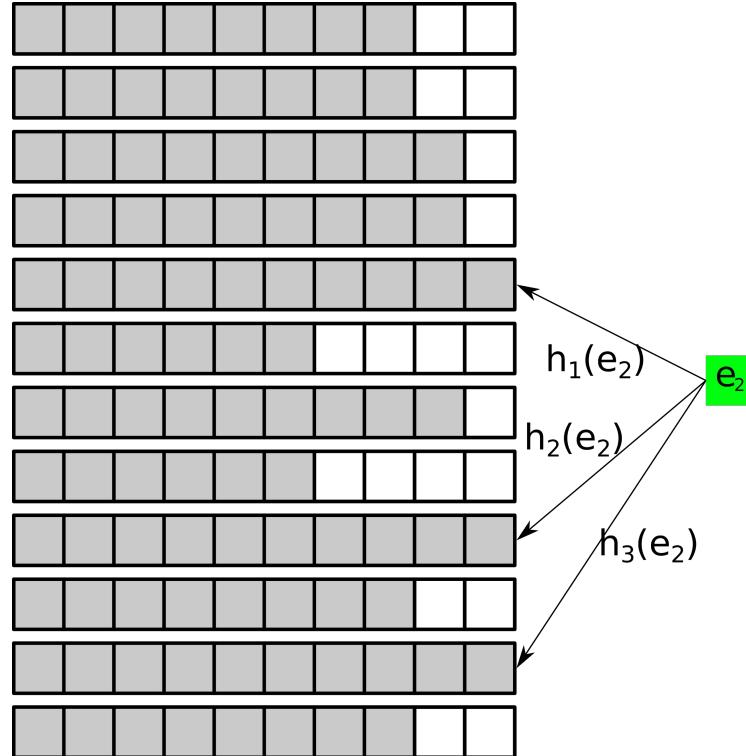
## (h,b) Cuckoo hashing: Random walk



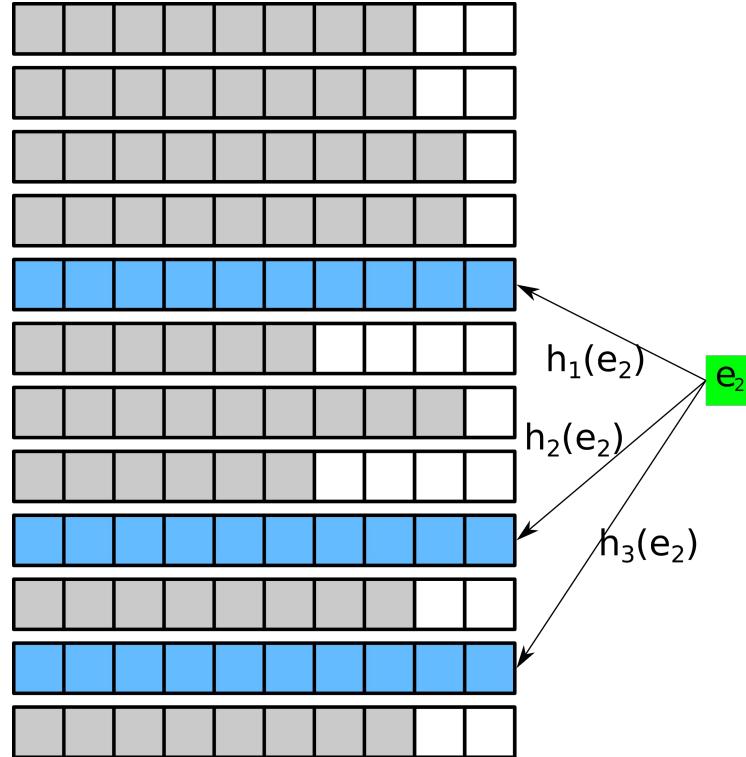
## (h,b) Cuckoo hashing: Random walk



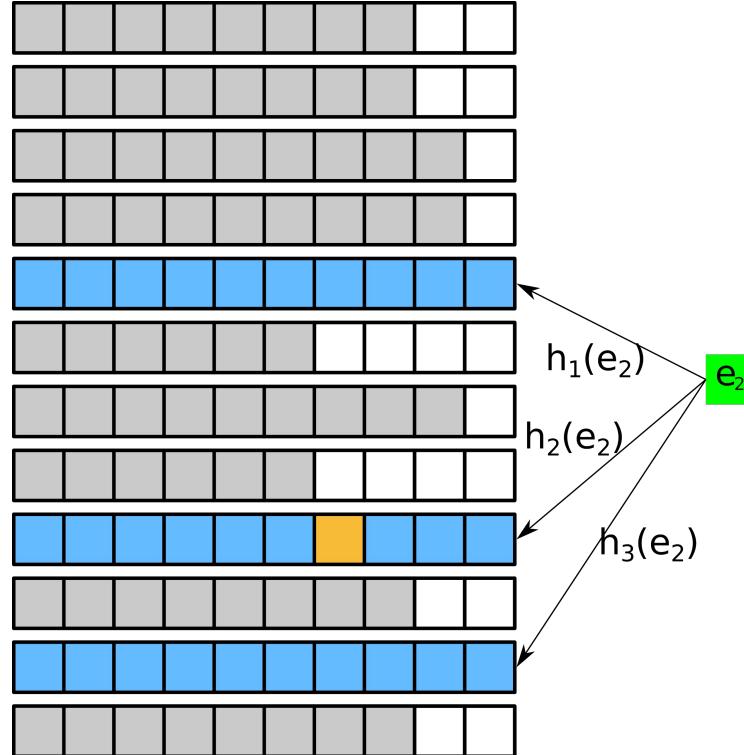
## (h,b) Cuckoo hashing: Random walk



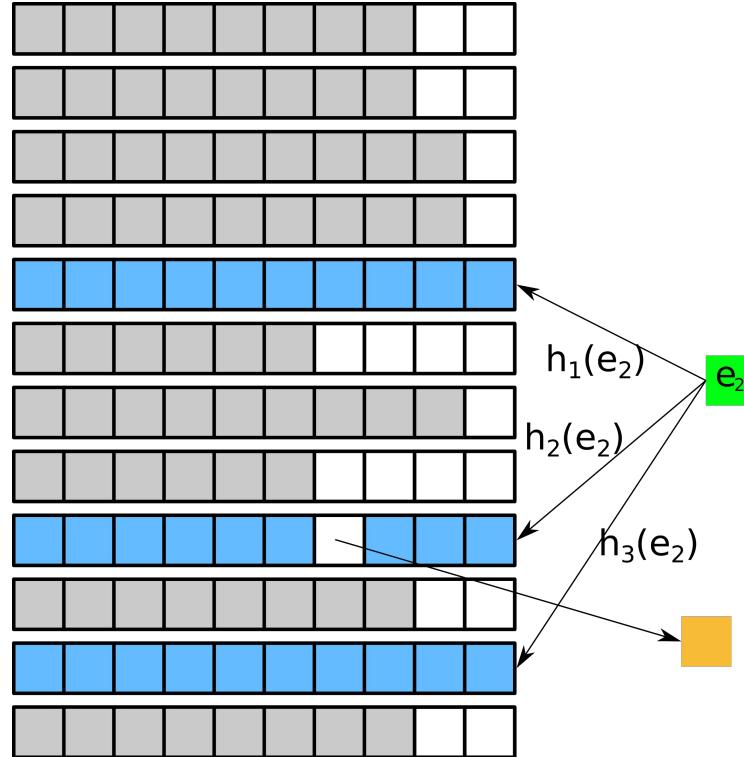
# (h,b) Cuckoo hashing: Random walk



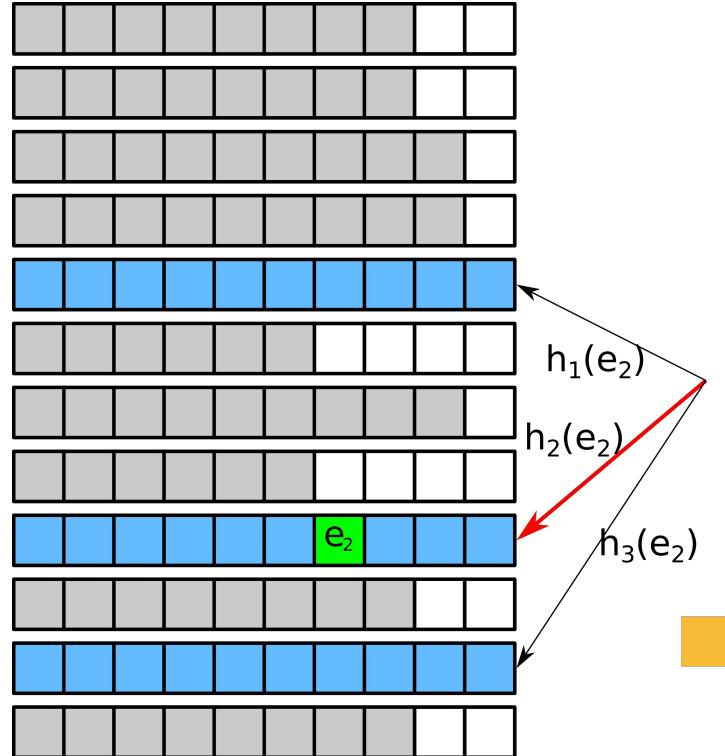
## (h,b) Cuckoo hashing: Random walk



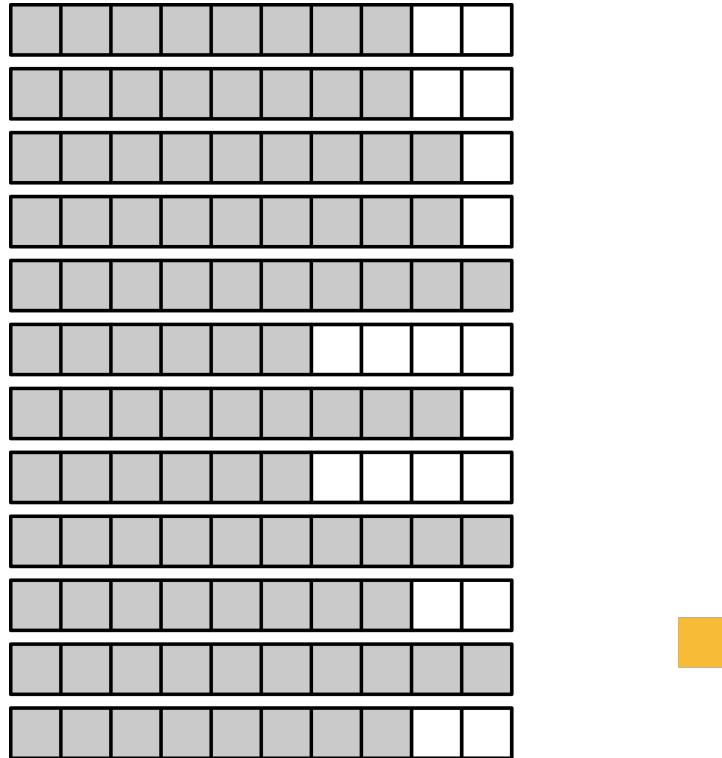
# (h,b) Cuckoo hashing: Random walk



## (h,b) Cuckoo hashing: Random walk



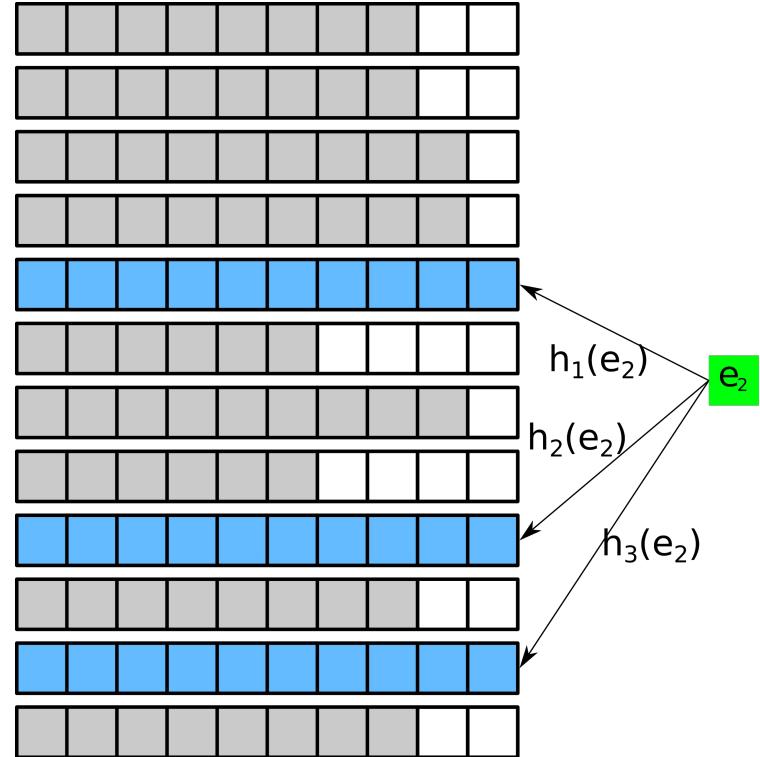
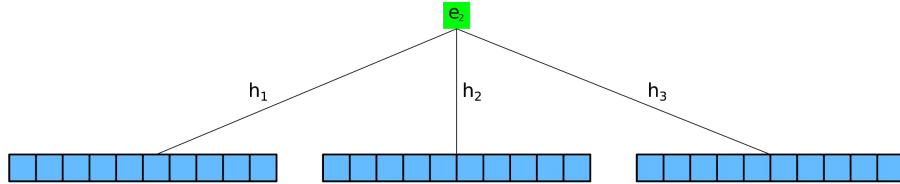
# (h,b) Cuckoo hashing: Random walk



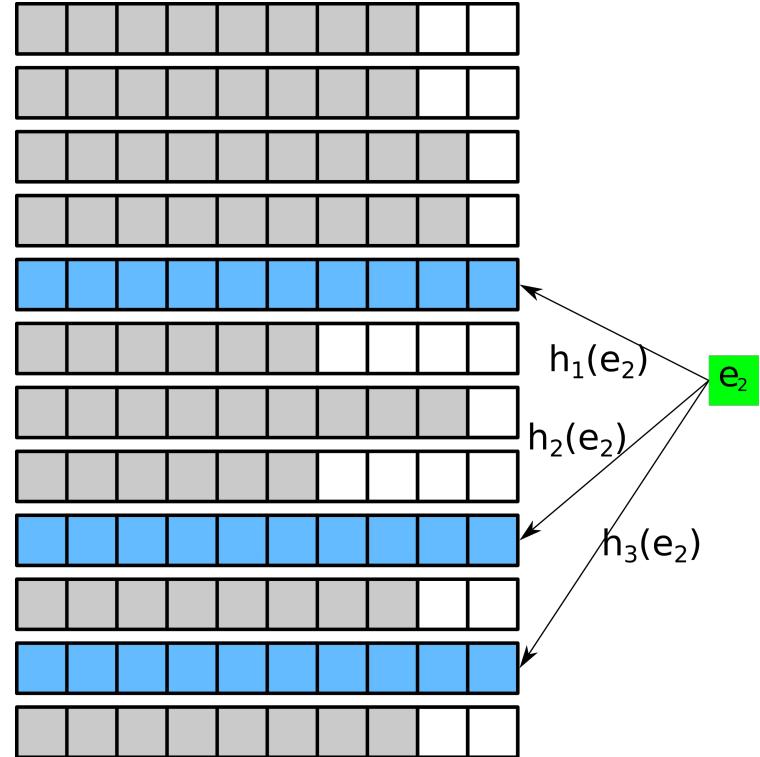
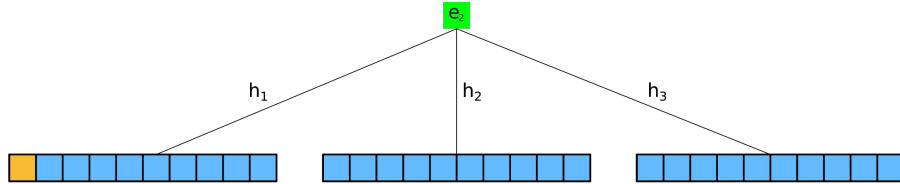
## (h,b) Cuckoo hashing: BFS

- BFS always finds a path to insert a new element (if one exists)
- BFS starts at the new elements
  - Check all reachable buckets
  - if there exists no empty slot:
    - Follow all elements in these buckets to their alternative locations

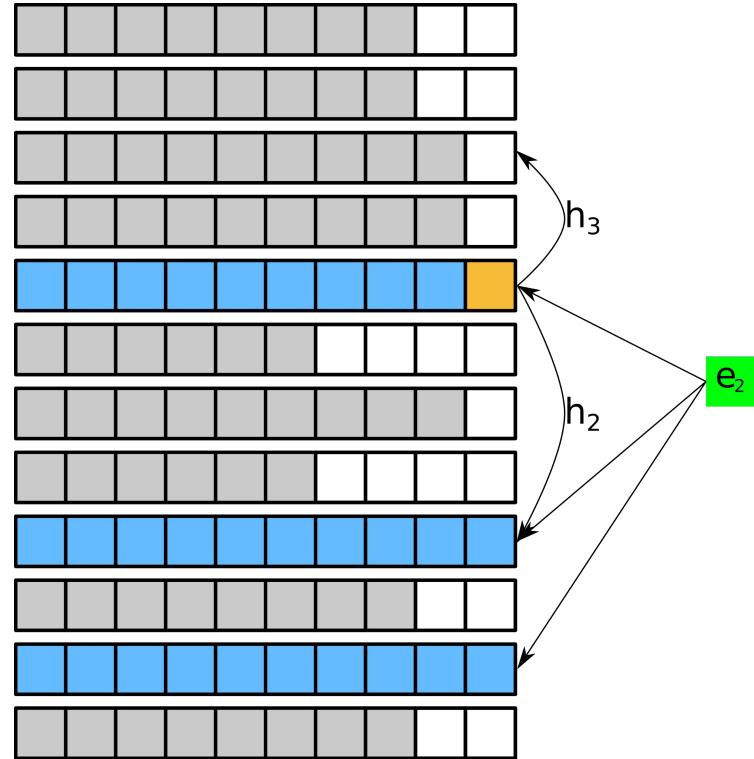
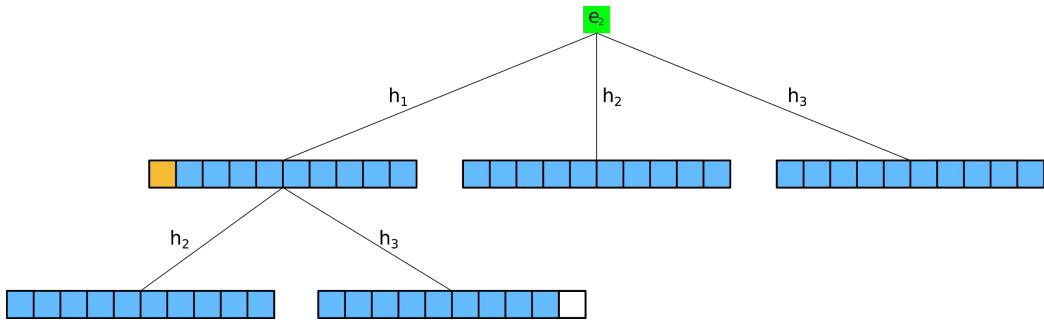
# (h,b) Cuckoo hashing: BFS



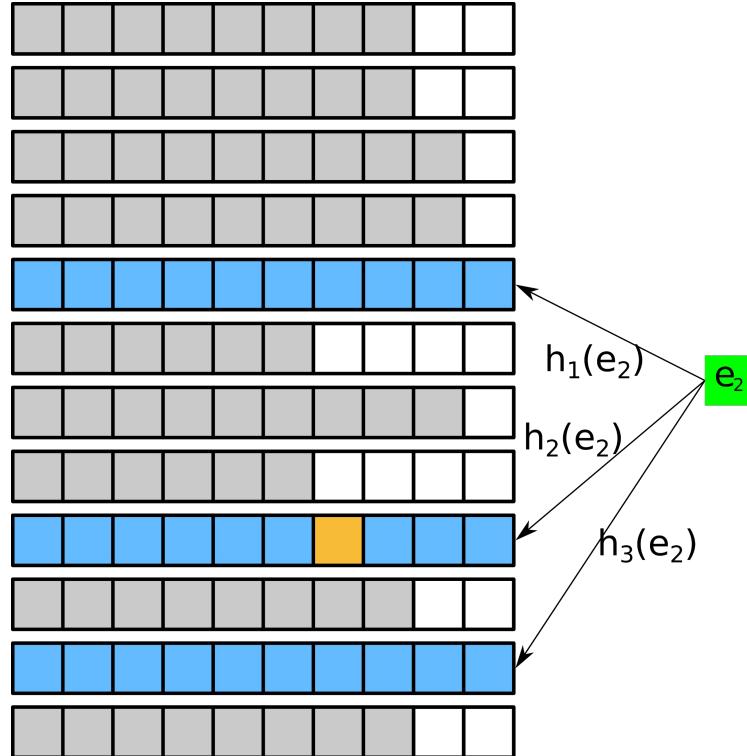
# (h,b) Cuckoo hashing: BFS



# (h,b) Cuckoo hashing: BFS



## (h,b) Cuckoo hashing: Query



# Achievable loads for (h,b) Cuckoo hashing

Classical (2,1) Cuckoo hashing: 0.5

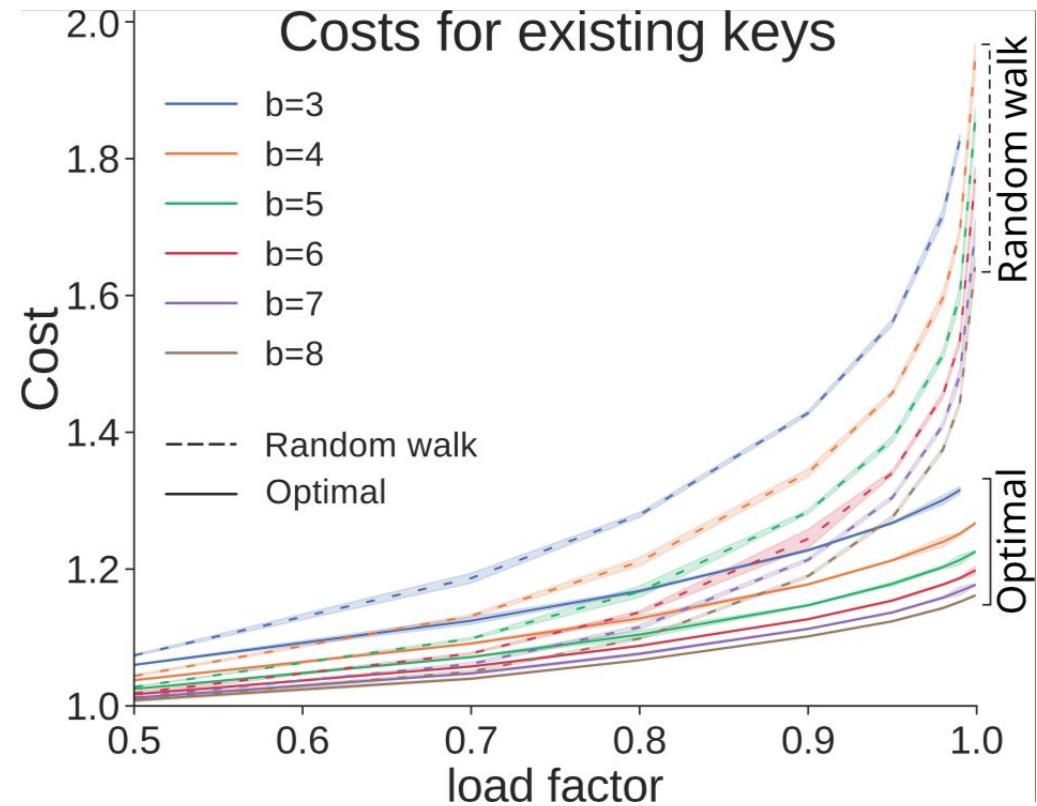
- Table from paper with load thresholds
- Expected lookups for random walk (Graphs) for (3, b)
- Werbung für Poster

# Achievable loads for (h,b) Cuckoo hashing

$b \ h$	2	3	4	5	6	7
1	0.5	0.9179352767	0.9767701649	0.9924383913	0.9973795528	0.9990637588
2	0.8970	0.9882014140	0.9982414840	0.9997243601	0.9999568737	0.9999933439
3	0.95915	0.9972857393	0.9997951434	0.9999851453	0.9999989795	0.9999999329
4	0.98037	0.9992531564	0.9999720661	0.9999990737	0.9999999721	0.9999999992

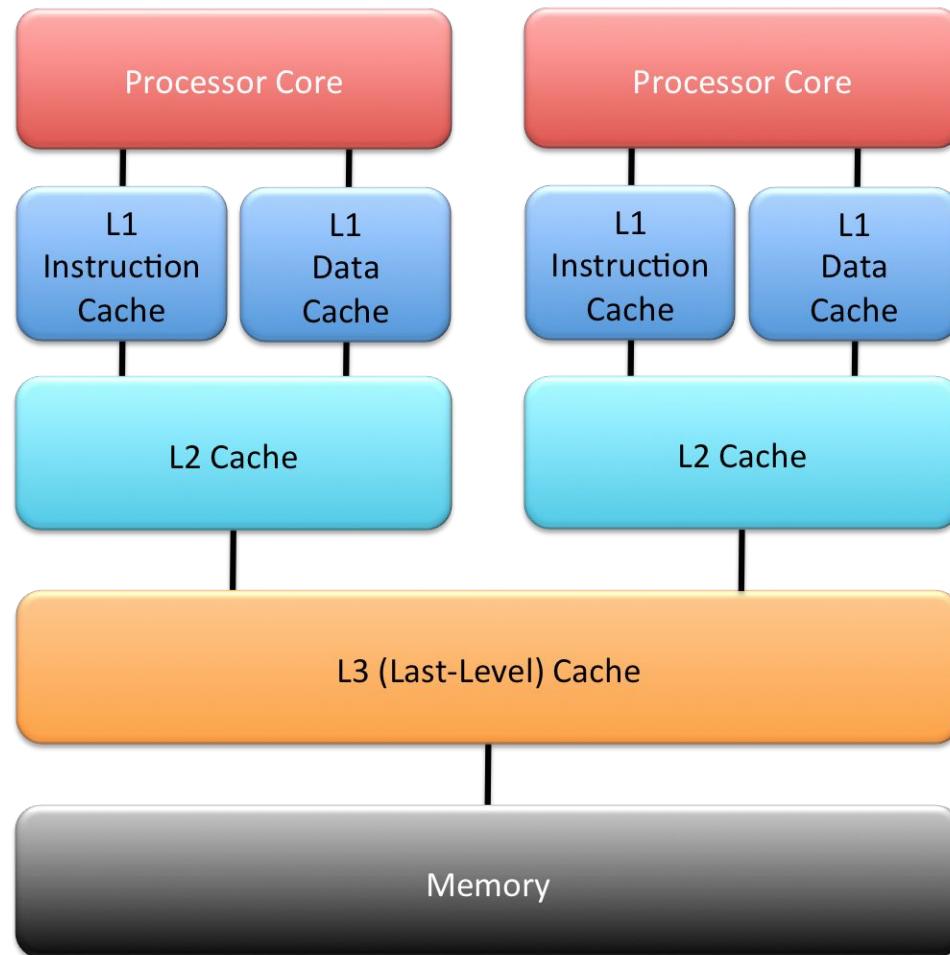
# Costs for $(h=3, b)$ Cuckoo hashing with random walk

- Bucket size  $b \in \{3, \dots, 8\}$
- Load factor  
 $\alpha \in \{0.5, 0.7, 0.8, 0.9, 0.95, 0.99, 0.999\}$
- Cost: expected number of memory lookups per lookup of a present key
- Even with  $h = 3$ , close to 1 lookup suffices on average (loads < 0.95). Worst case of 3 happens rarely.

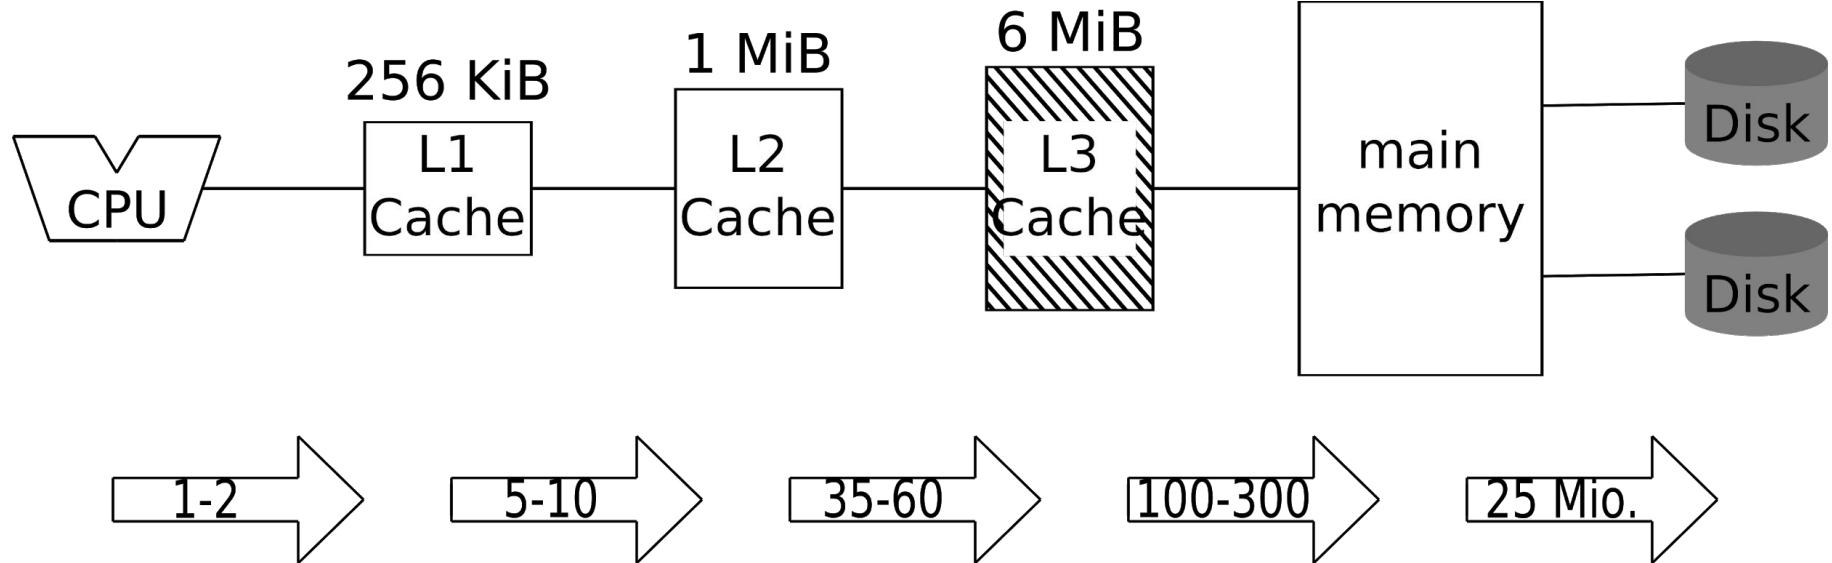


# The role of modern hardware

# CPU caches

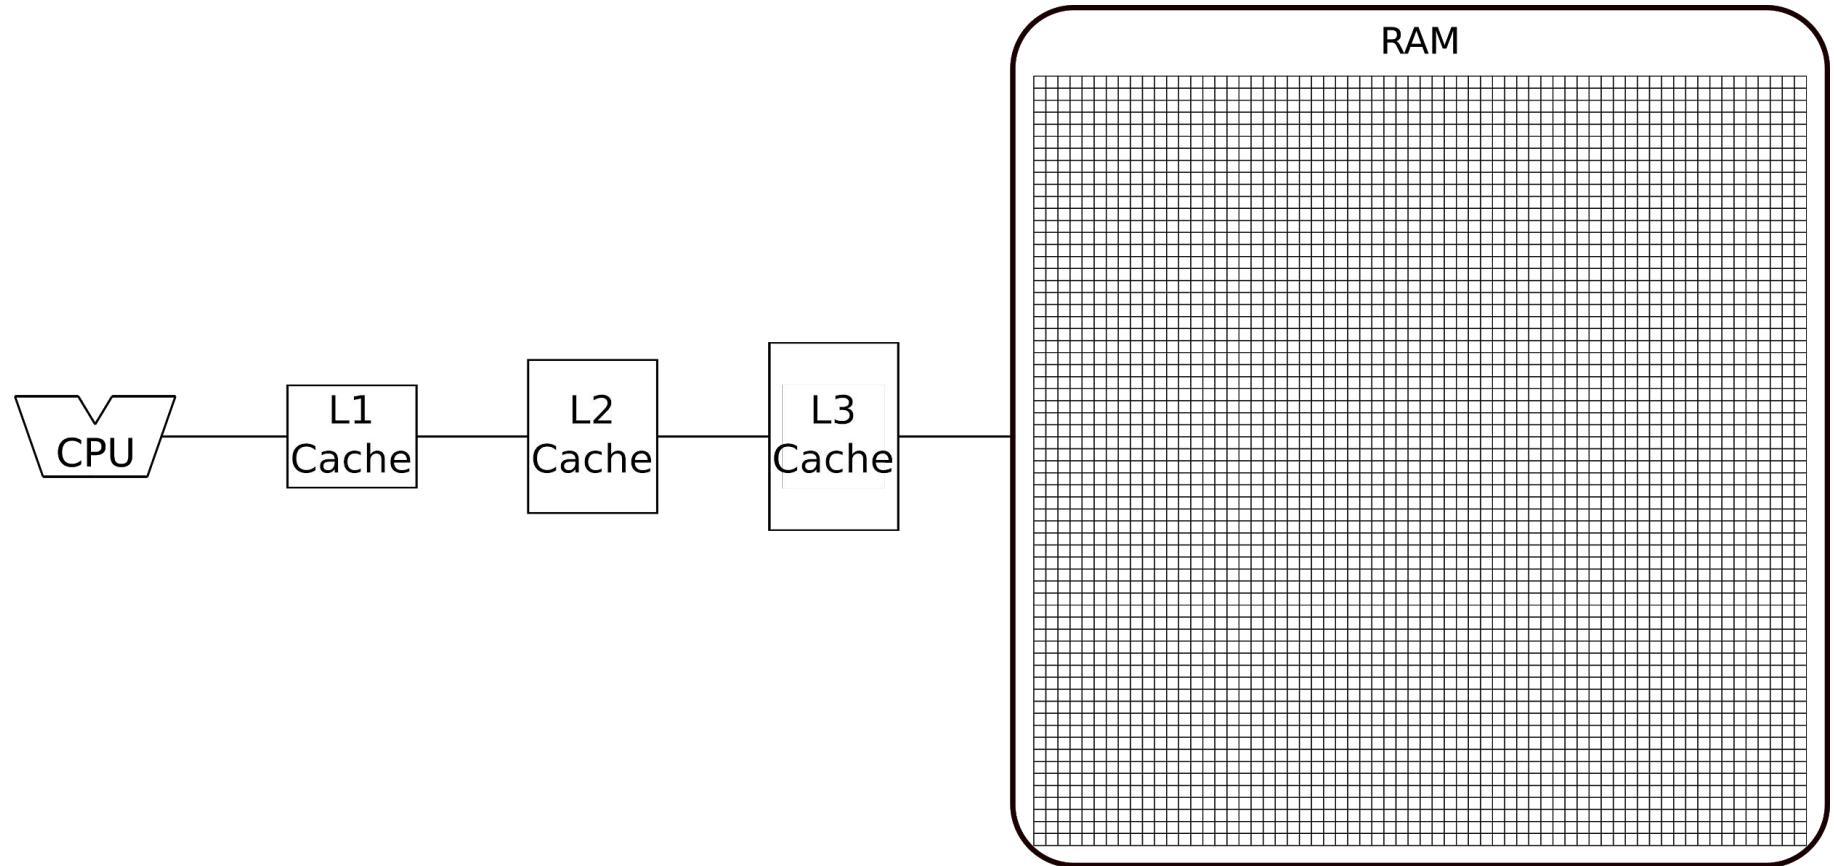


# CPU caches

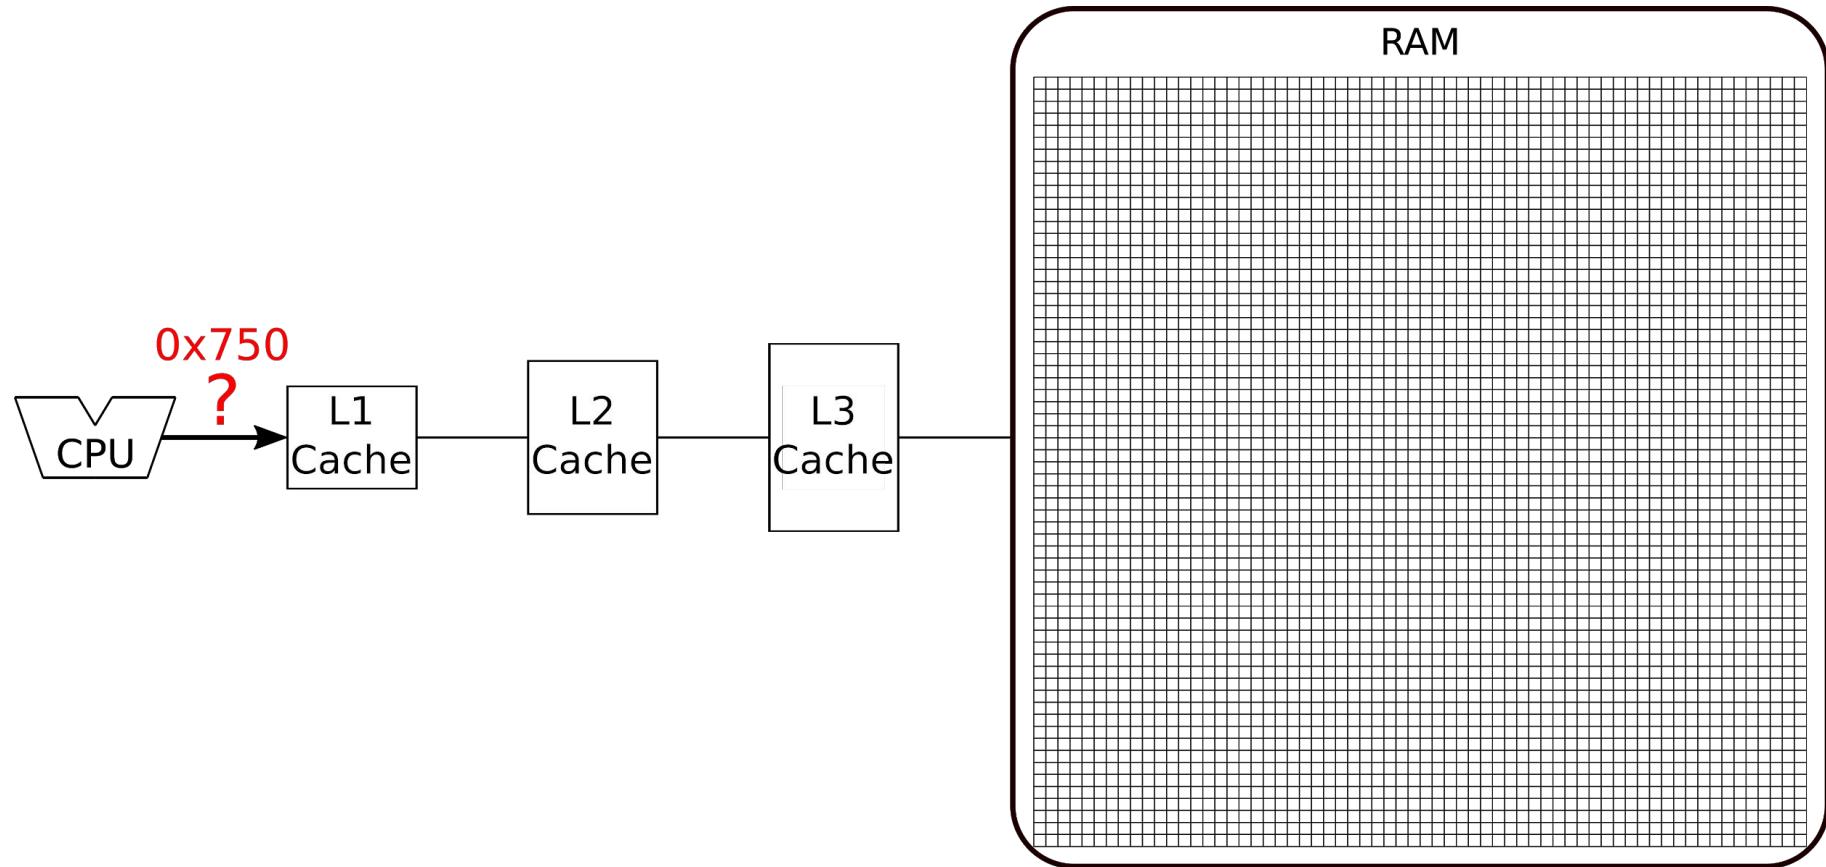


This and following illustrations by Uriel Elias Wiebelitz (TU Dortmund)

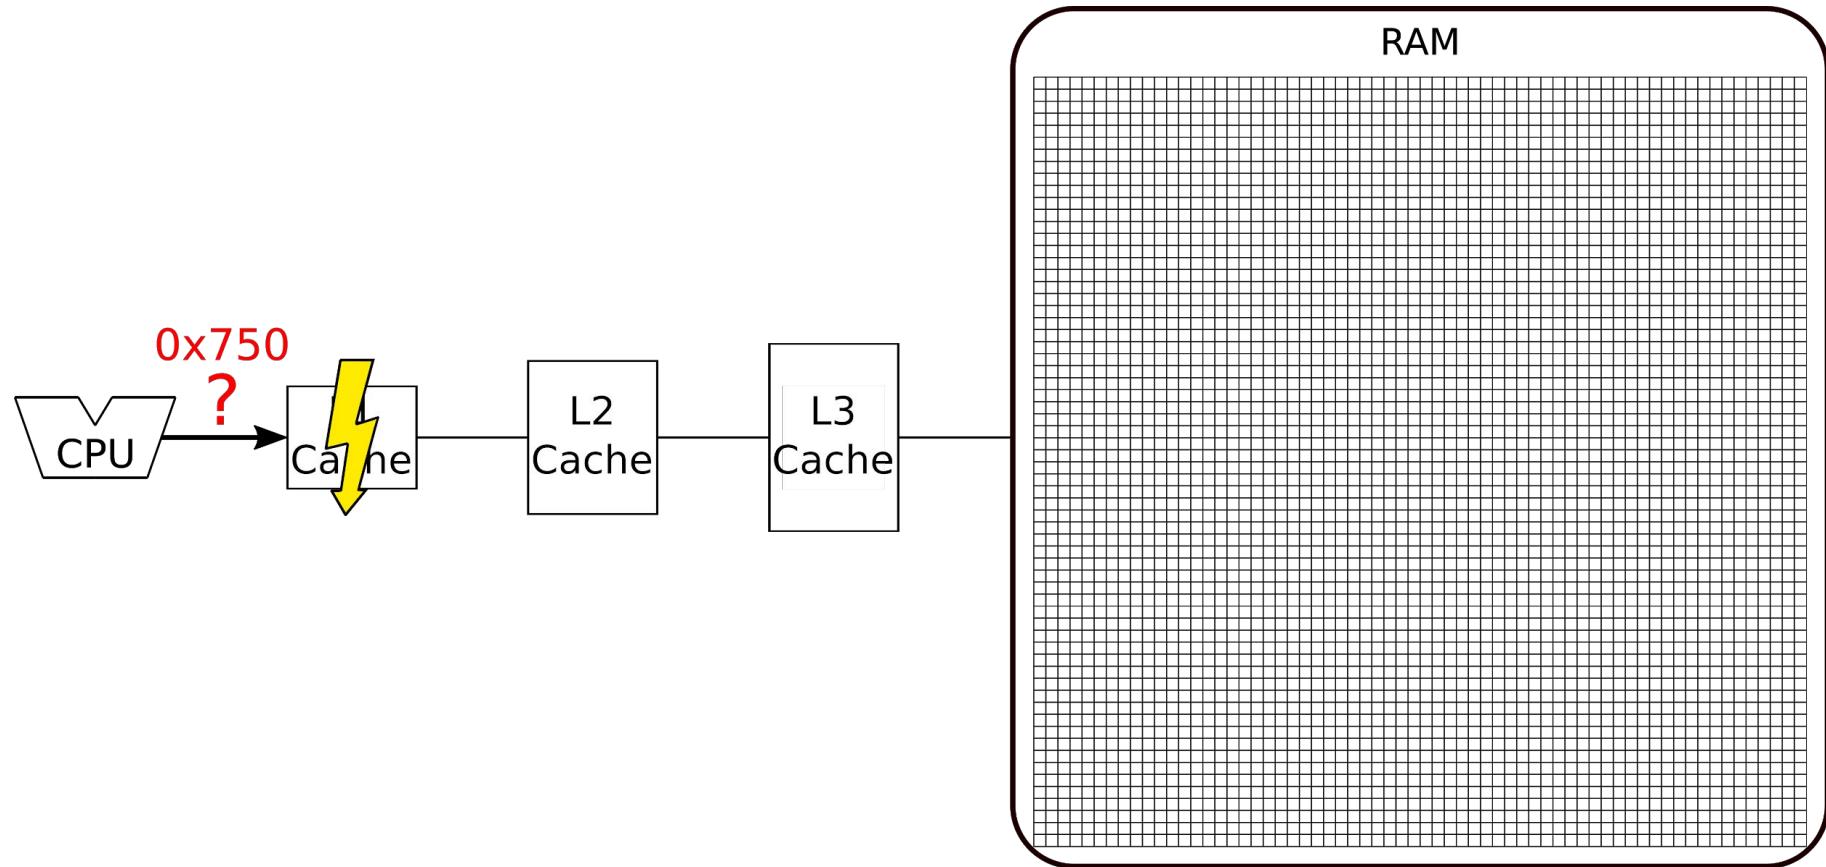
# Cache line



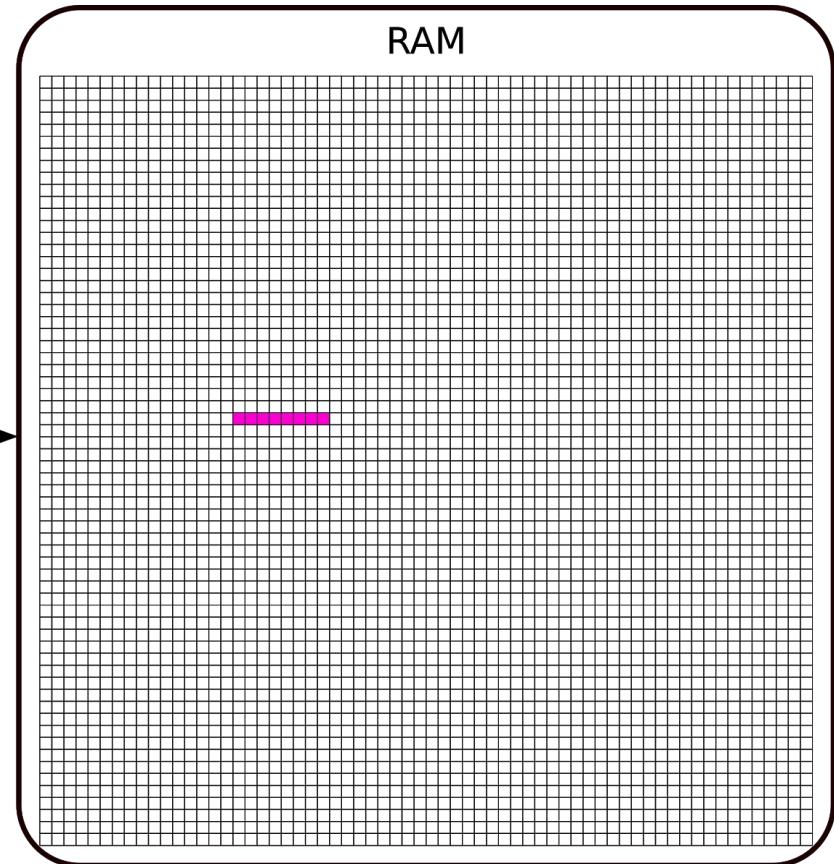
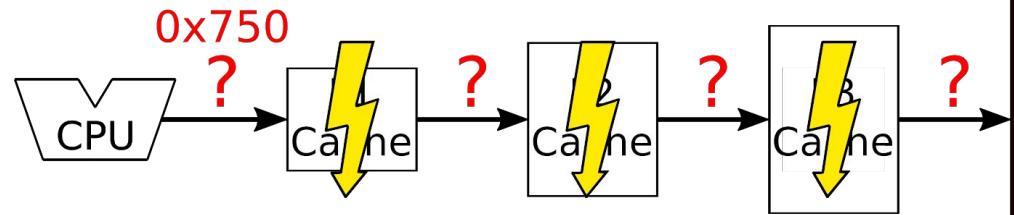
# Cache line



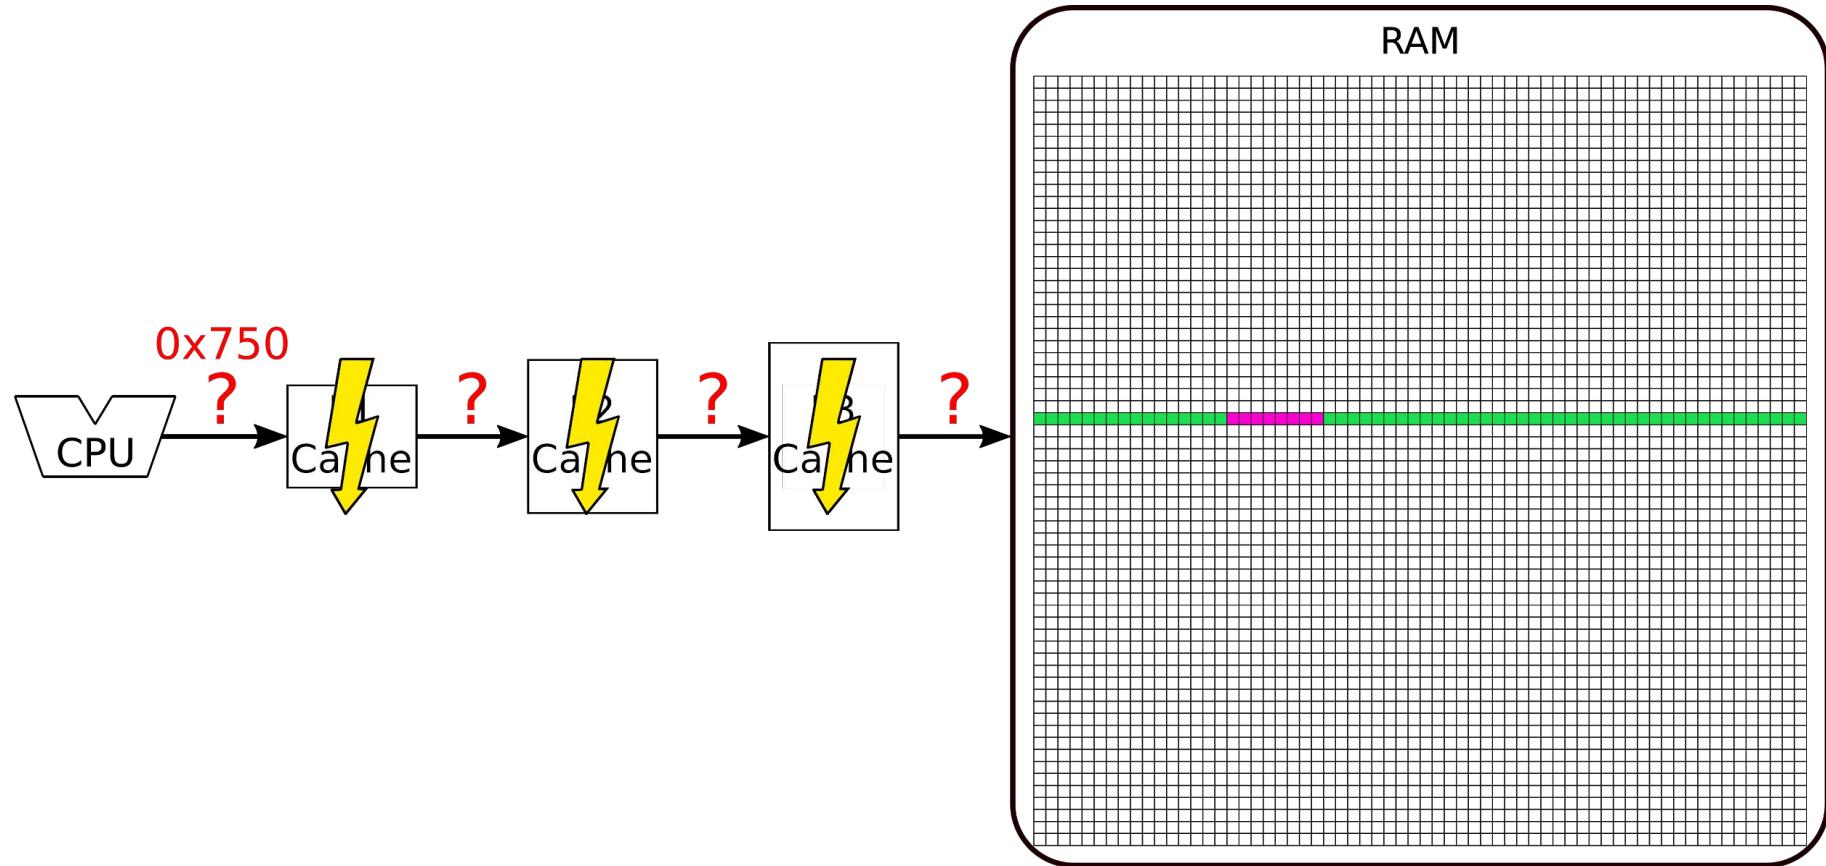
# Cache line



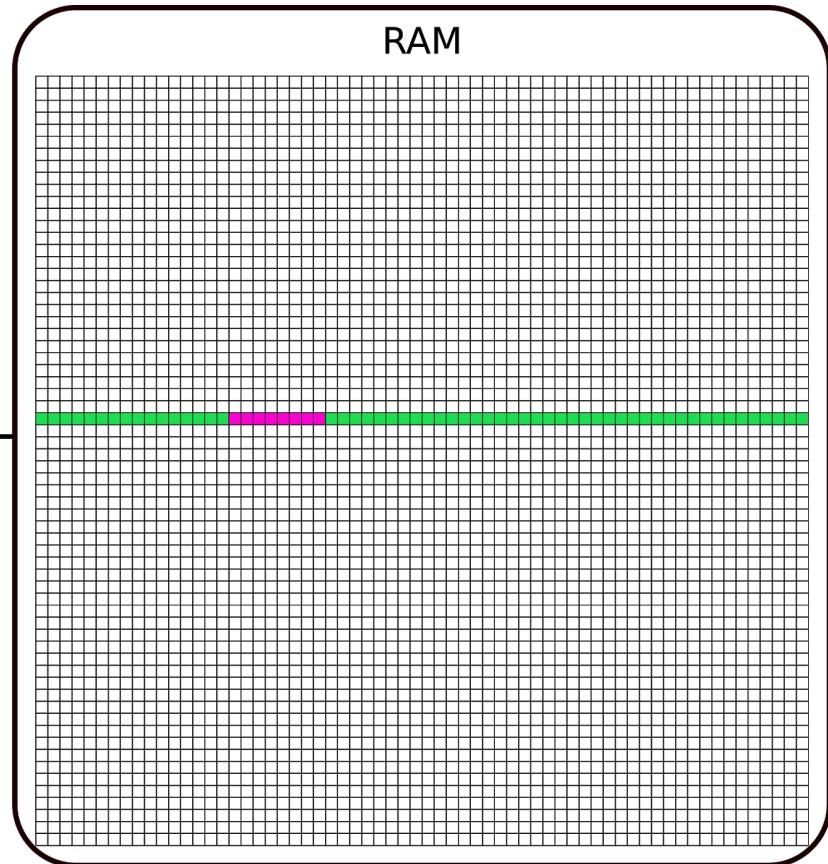
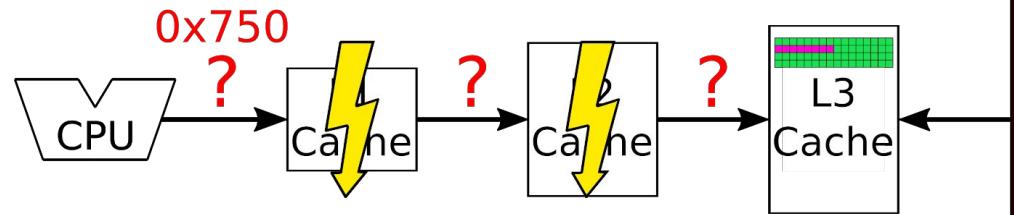
# Cache line



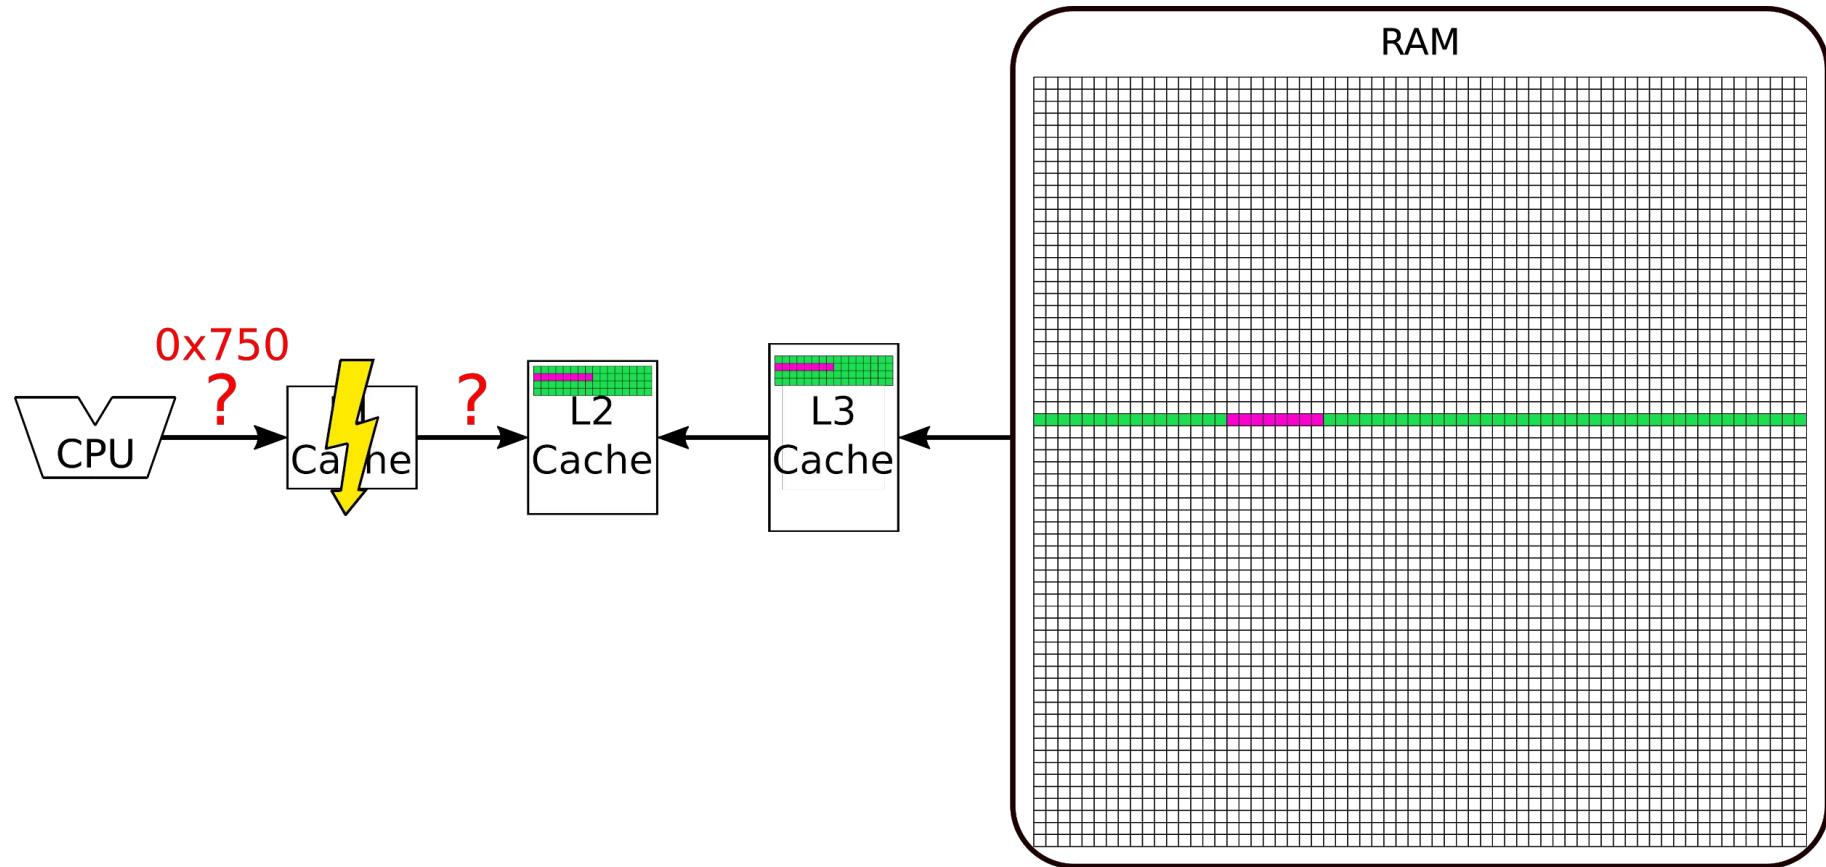
# Cache line



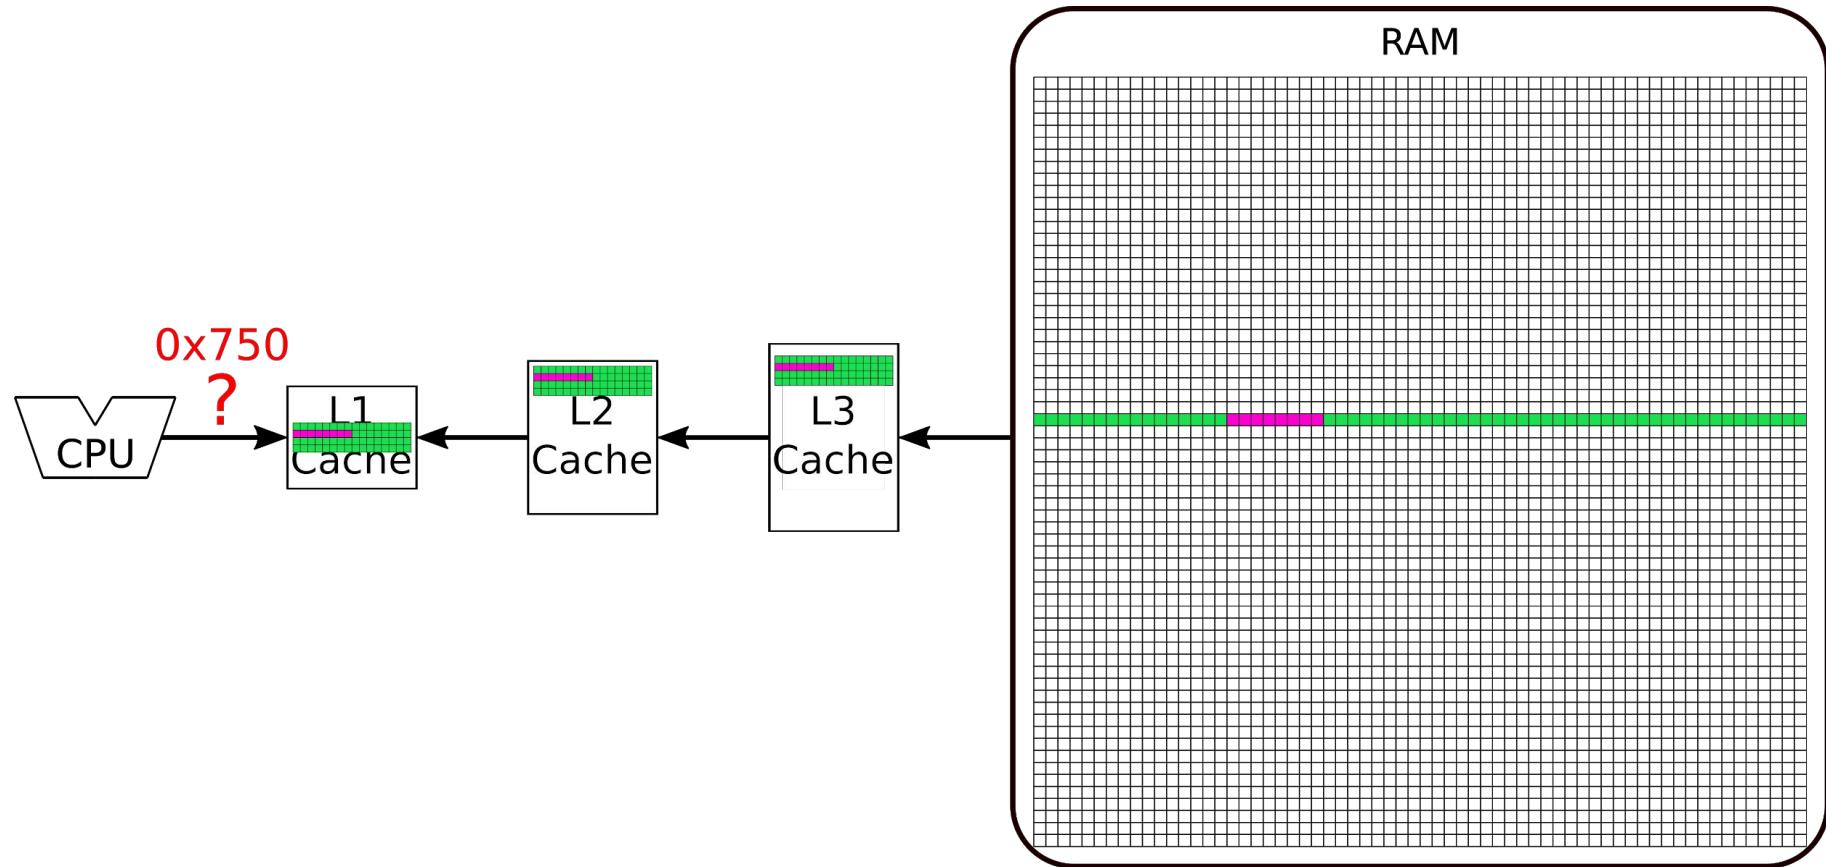
# Cache line



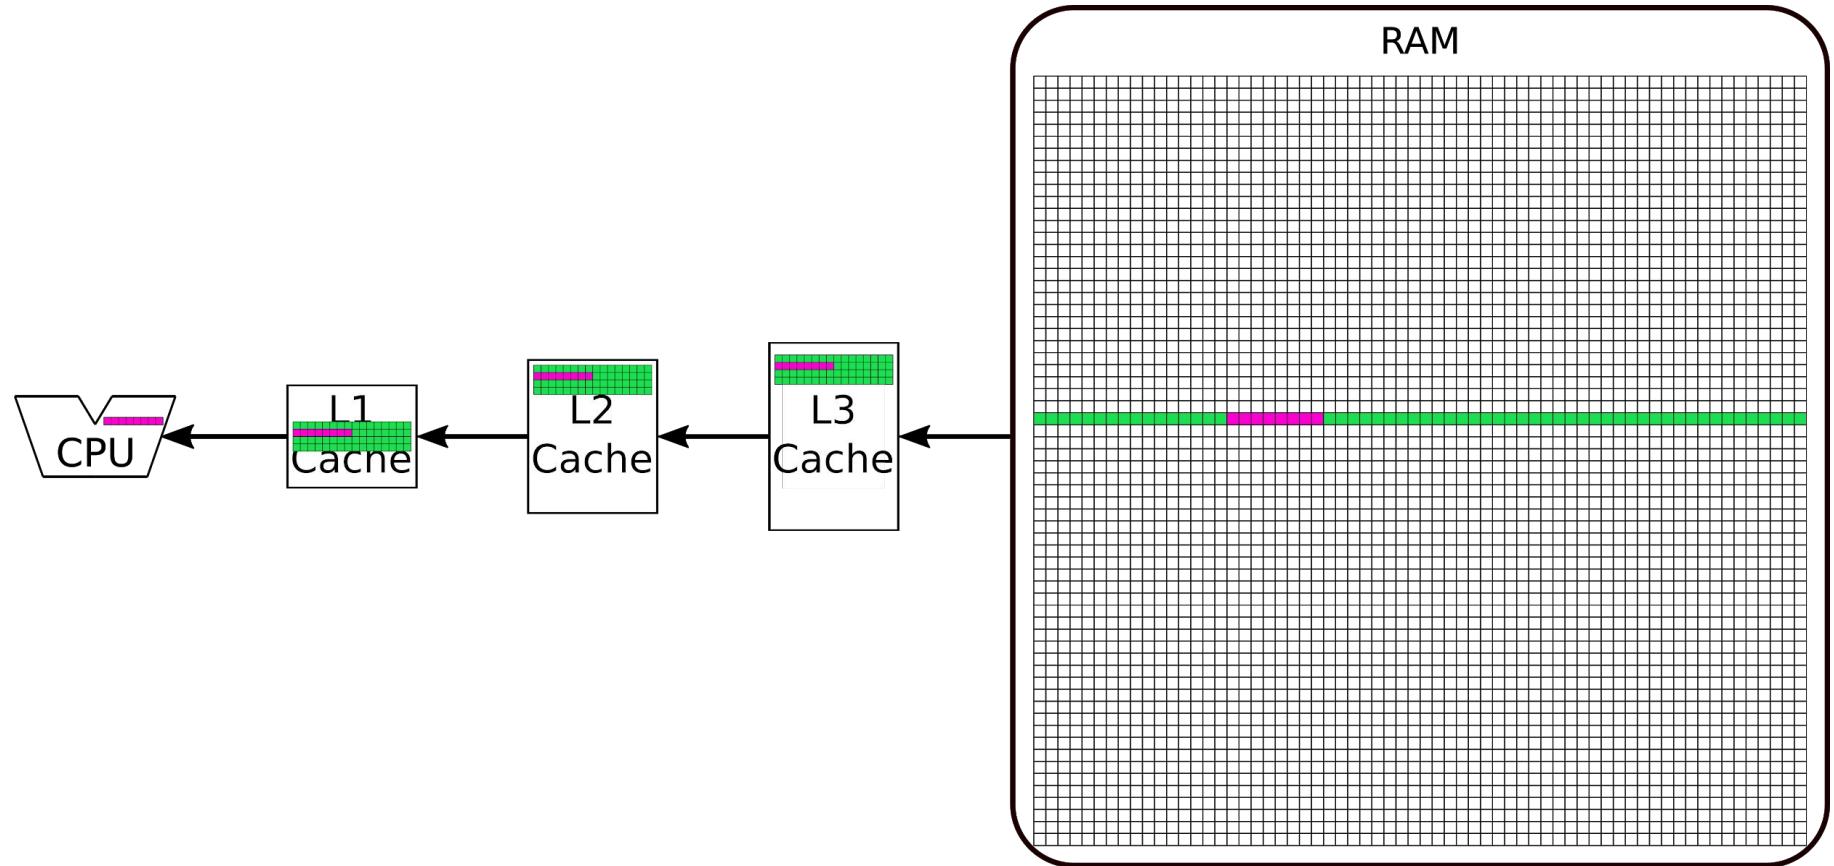
# Cache line



# Cache line



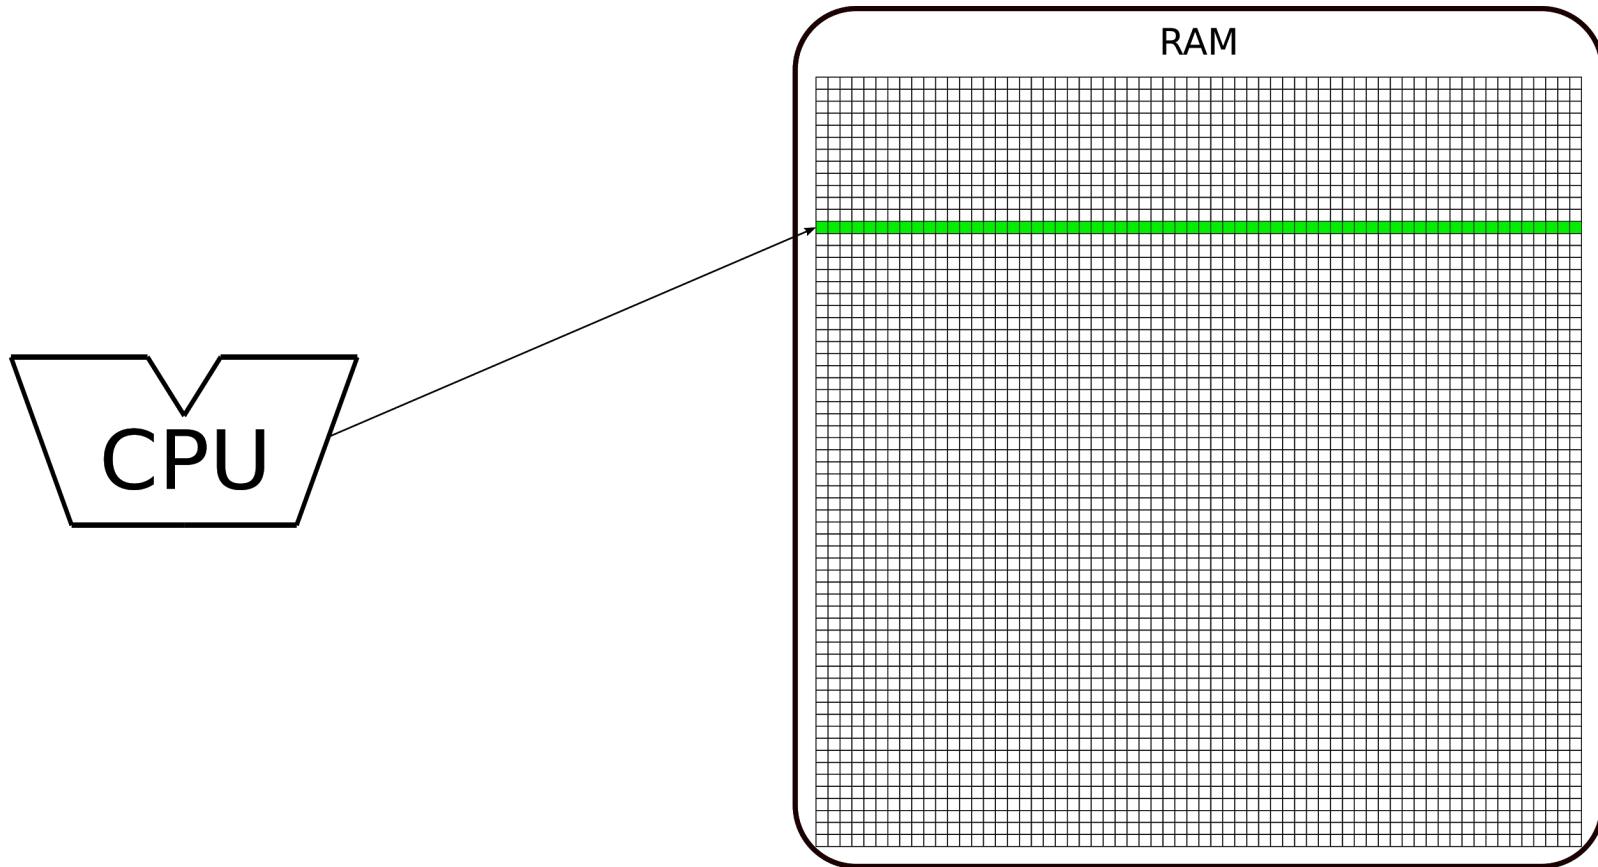
# Cache line



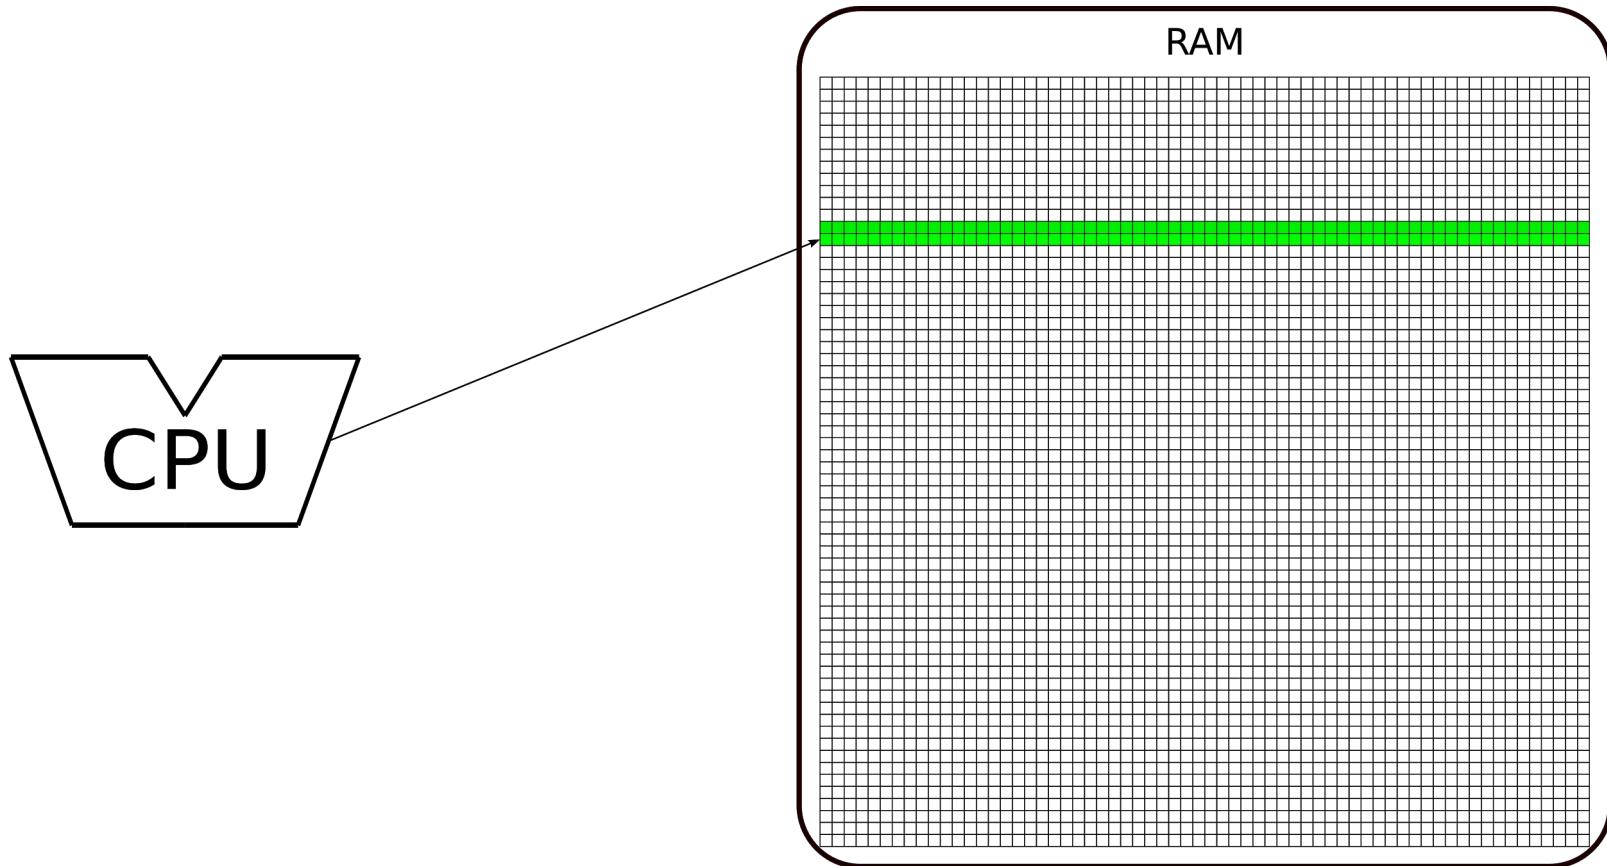
# Prefetching

- Reduce the waiting period for missing data
- Easy access patterns are prefetched by the hardware
  - Linear access in both directions
  - Jumps of fixed width
- Complex patterns need manual prefetching (software prefetching)

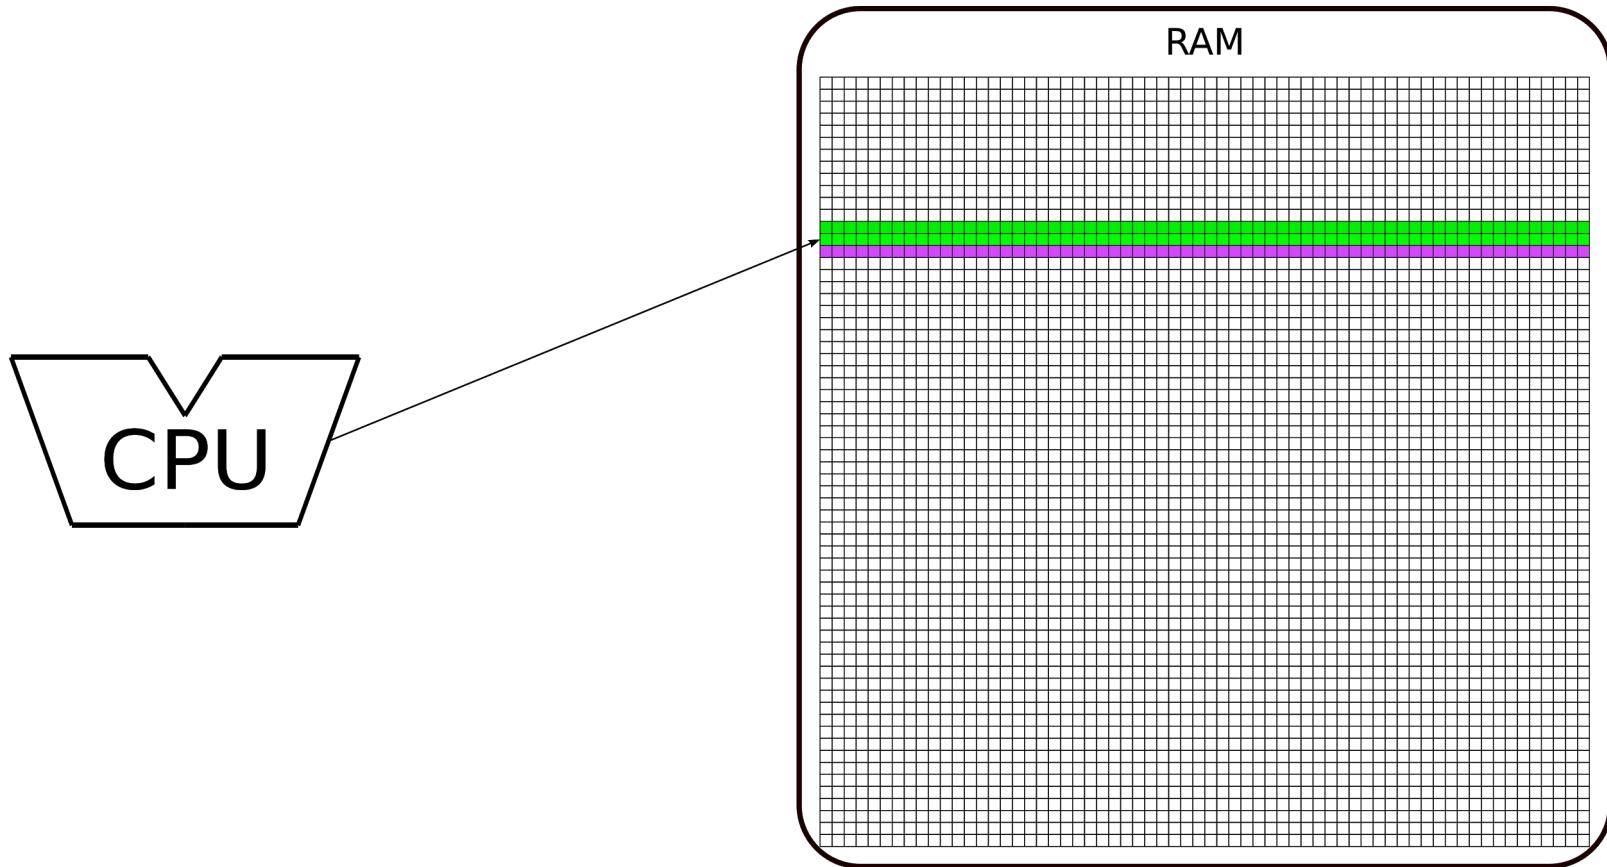
# Hardware prefetching



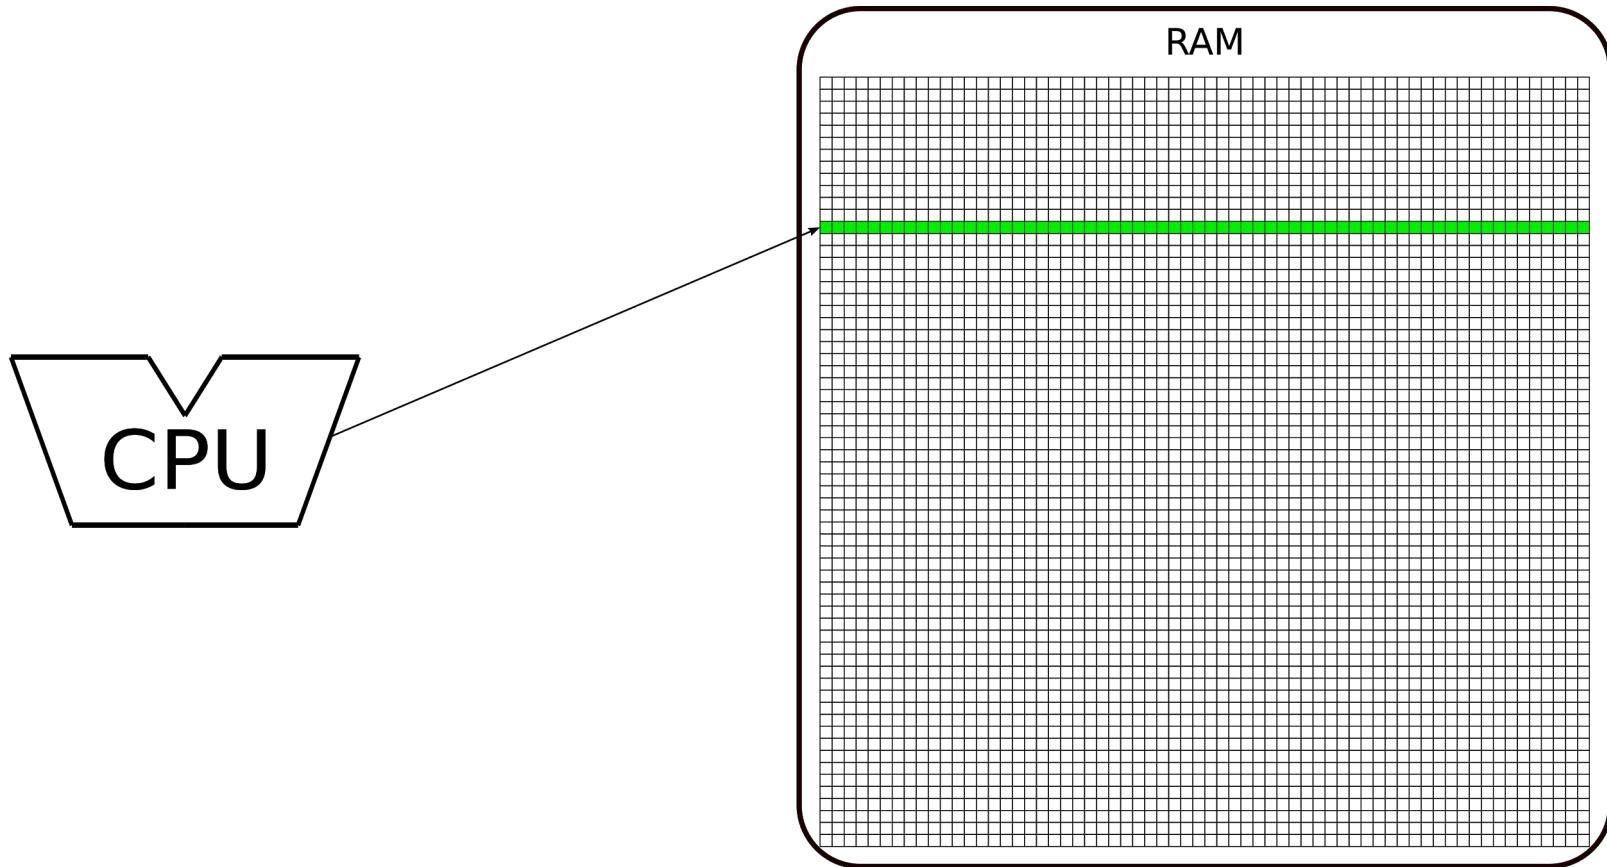
# Hardware prefetching



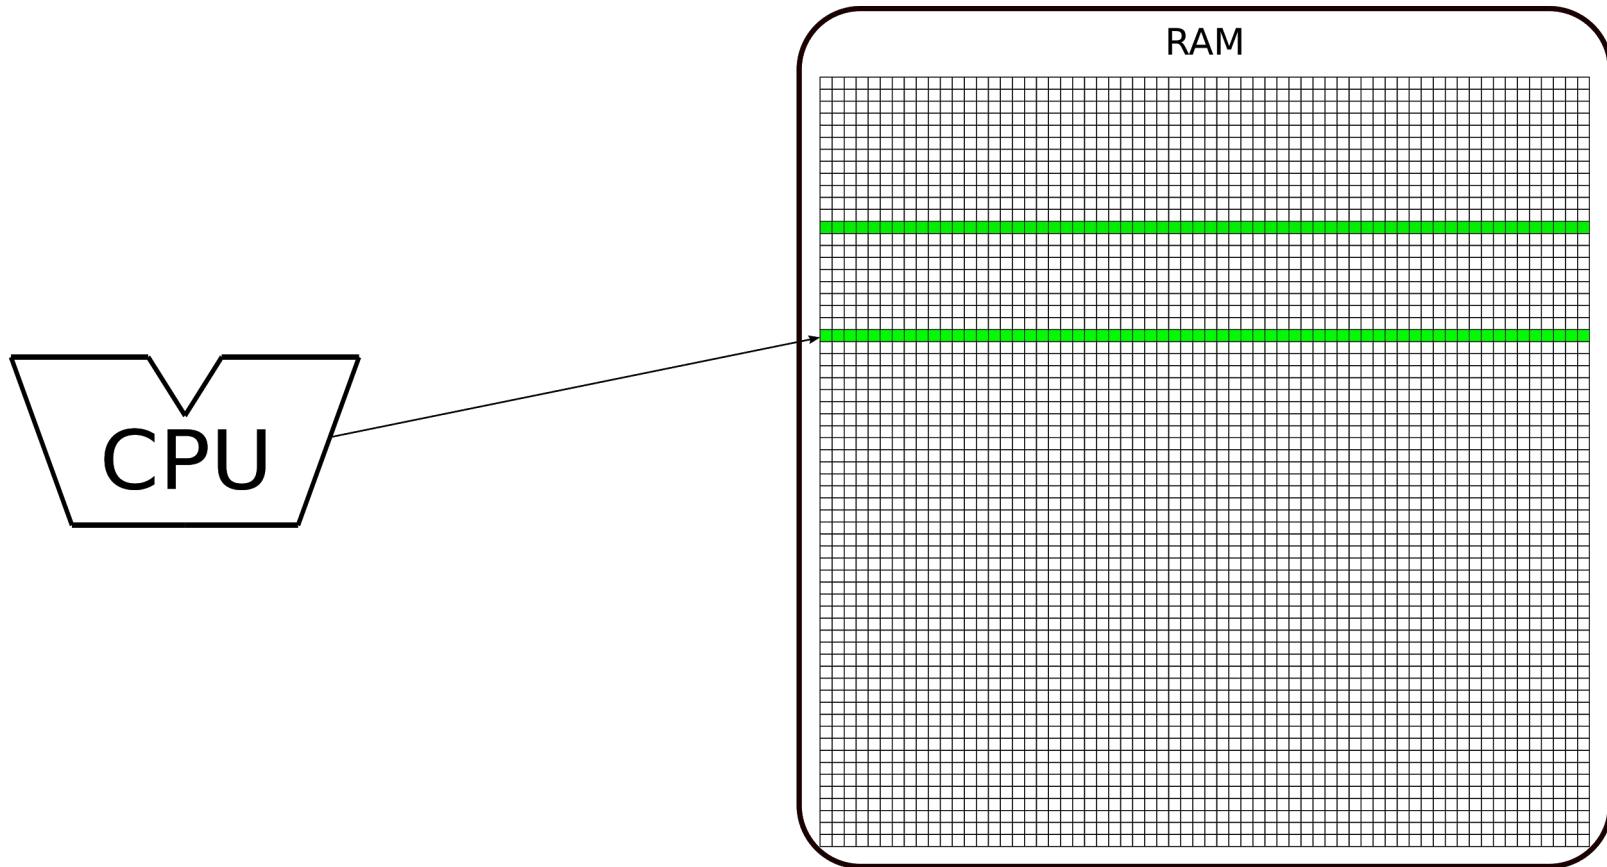
# Hardware prefetching



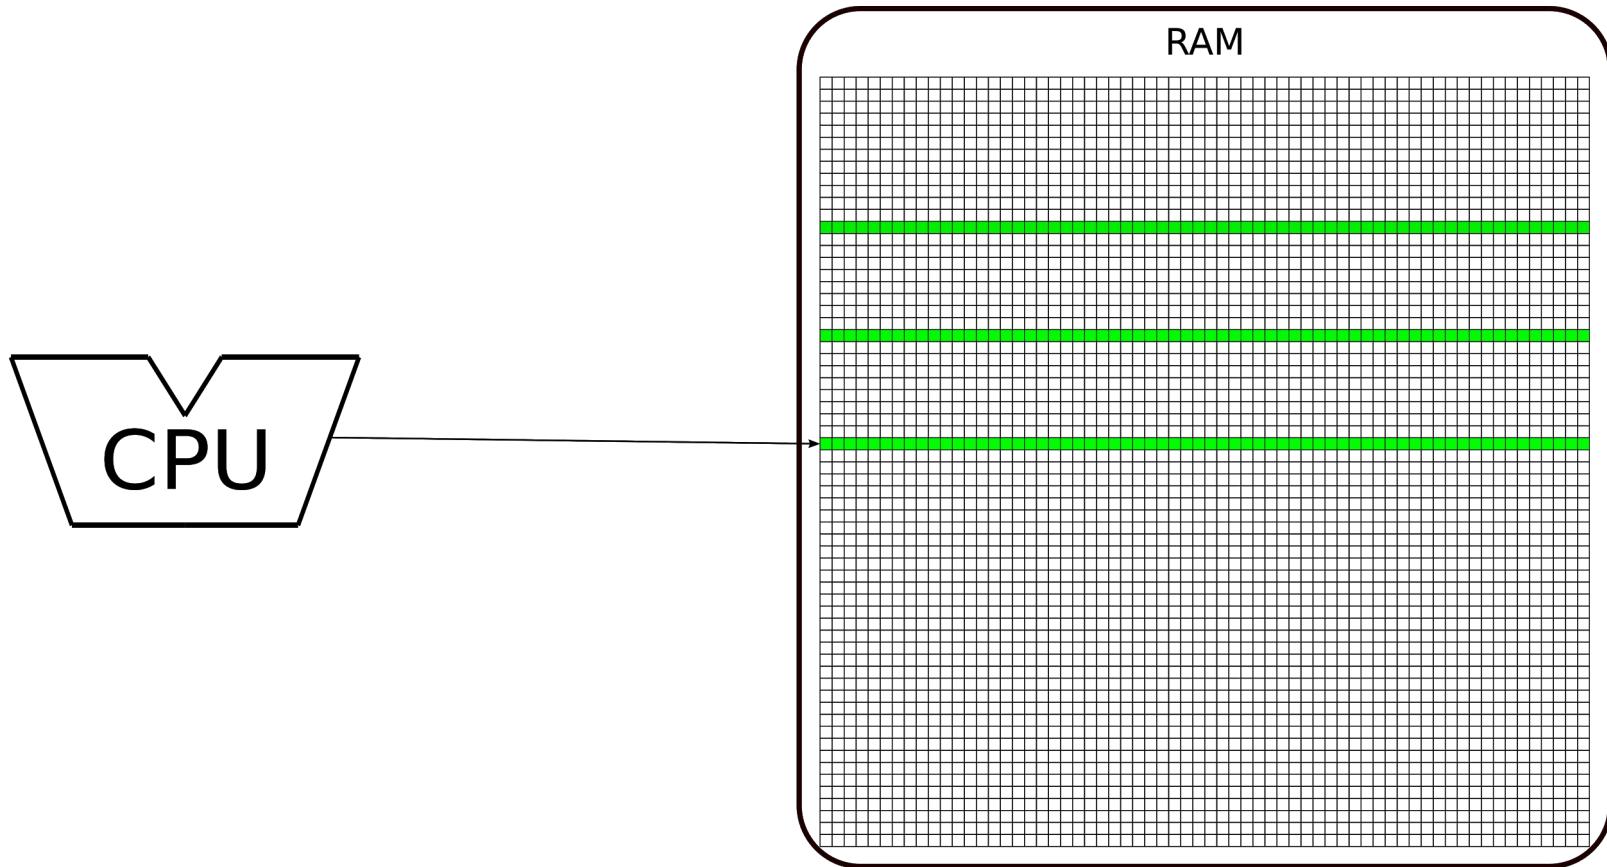
# Hardware prefetching



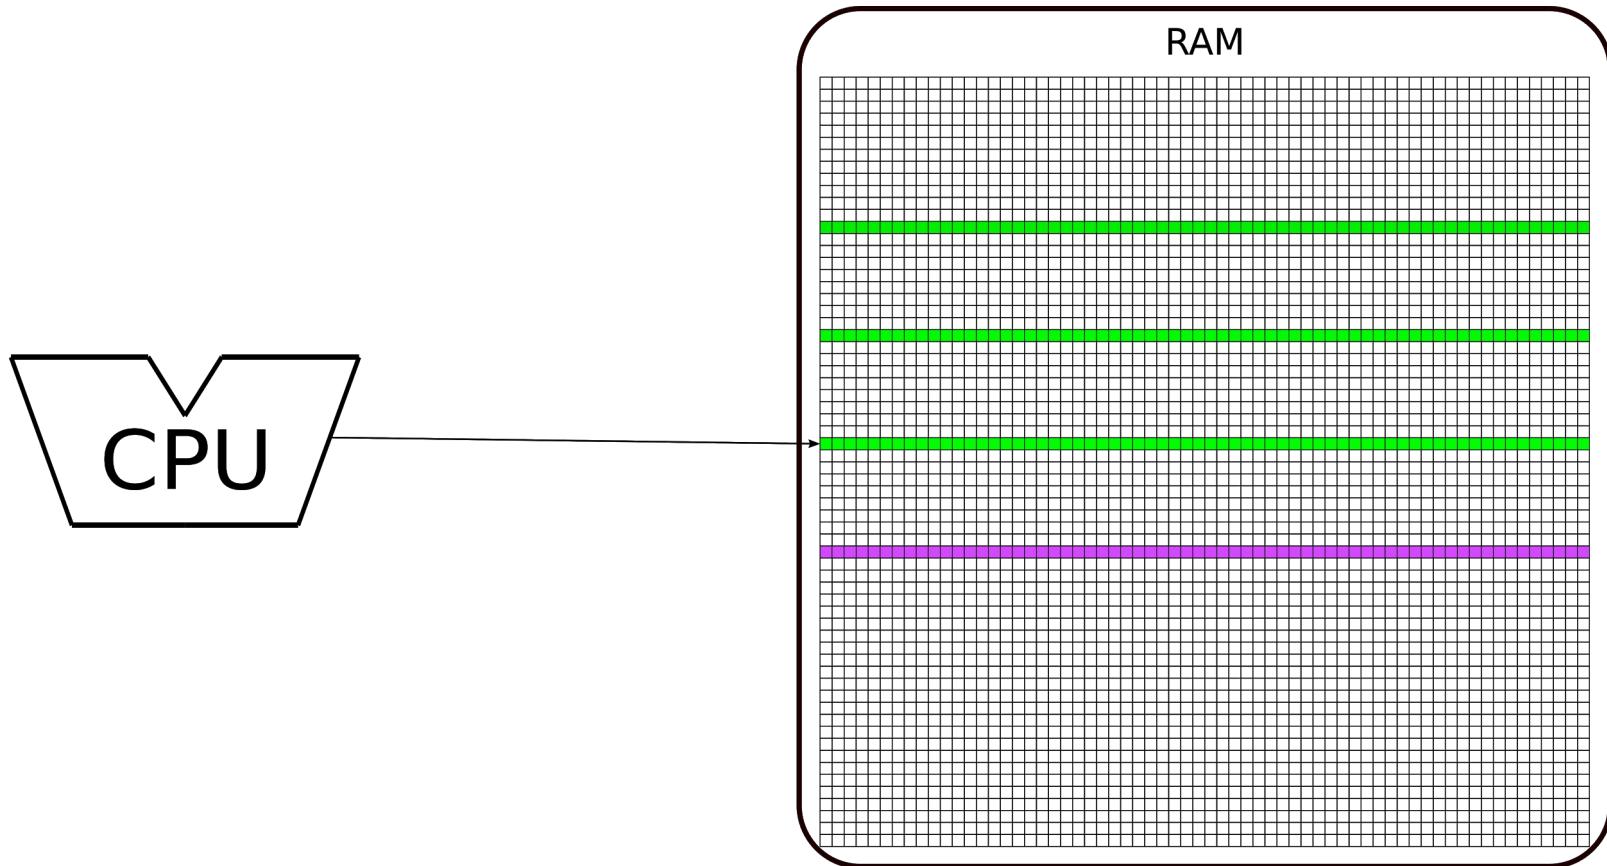
# Hardware prefetching



# Hardware prefetching

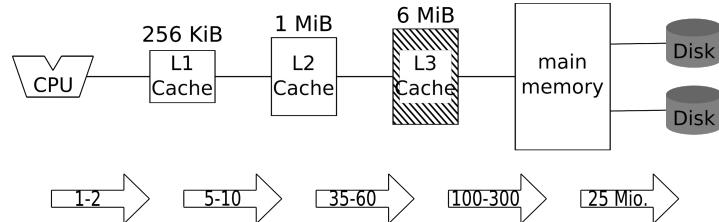


# Hardware prefetching



# Efficiency of linear probing

- Cache hierarchy + hardware prefetching for regular access patterns means:
- Hopscotch:
  - check consecutive positions → cache-efficient
- Cuckoo hashing:
  - Searching within a bucket is cache-efficient
  - Looking up a bucket is not, but limited to  $h=3$  buckets.
  - Also, there is software prefetching !



# Software prefetching

- CPU instruction
- Can slow down the program
  - One instruction more to handle by the CPU
  - Still needed data can be removed from the cache
    - Needs to be loaded again

```
for(int i=0; i<1000; ++i) {  
    __builtin_prefetch(&arr[i + k]);  
    ++arr[i];  
}
```

# Prefetching in hash tables

- Strategies are based on multiple facts:
  - cache size
  - distribution of elements
  - size of needed data
- Hard to predict effects theoretically,  
needs benchmarking !

# Parallelism

So far only serial algorithms, modern hardware is multi-core, so can we parallelize

- a) insertion
- b) lookup?

b) is easy (read-only) but a) is difficult (data races)

Synchronization needed (locks, spin locks ?)

CAS instruction (Jellyfish)

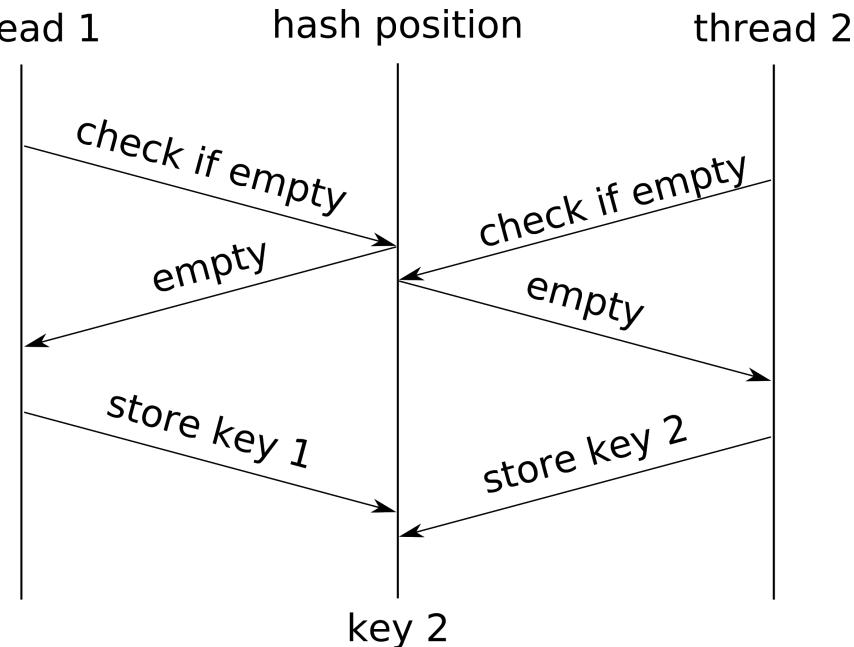
SIMD instructions (single instructions multiple data)

# Parallelism

- So far only serial algorithms
- Modern hardware is multi-core
- Also SIMD: single instruction multiple data  
(e.g. compute hash functions on multiple  $k$ -mers in parallel)
- Parallel lookup is easy:
  - data does not change
  - only read access
- Parallel write is harder:
  - Ensure that the data is always consistent
  - Multiple threads write to the same memory location:  
Synchronisation needed

# Access without synchronisation

- Both threads check whether the hash position is empty or not.
- Both see that the location is empty.
- Thread one stores key 1.
- Thread 2 stores key 2 and **overwrites** key 1.
- Key 1 is lost.

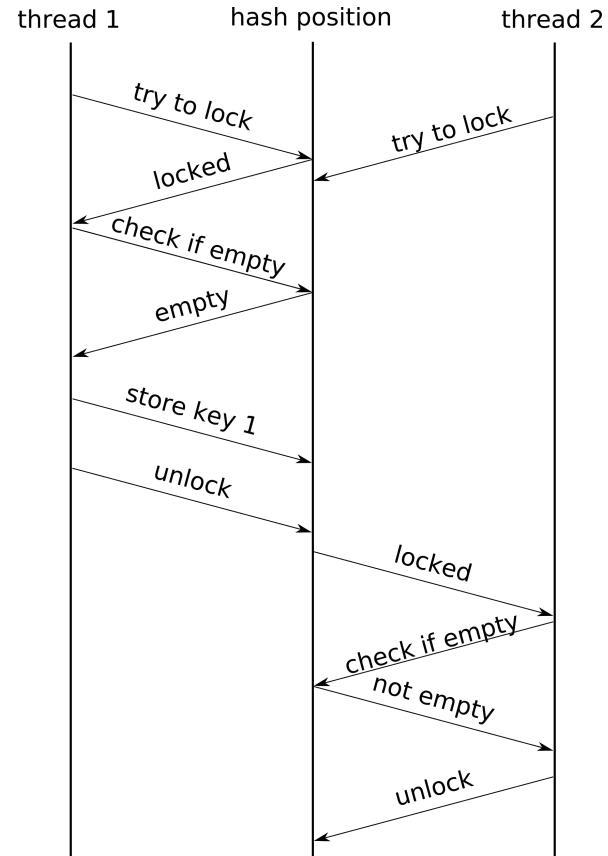


# Access with synchronisation

- Try to lock table slot
- As soon as the lock is confirmed:
  - change value in slot
- If slot is locked:
  - Wait until the lock can be obtained

Large memory overhead if explicit locks are used for every single slot.

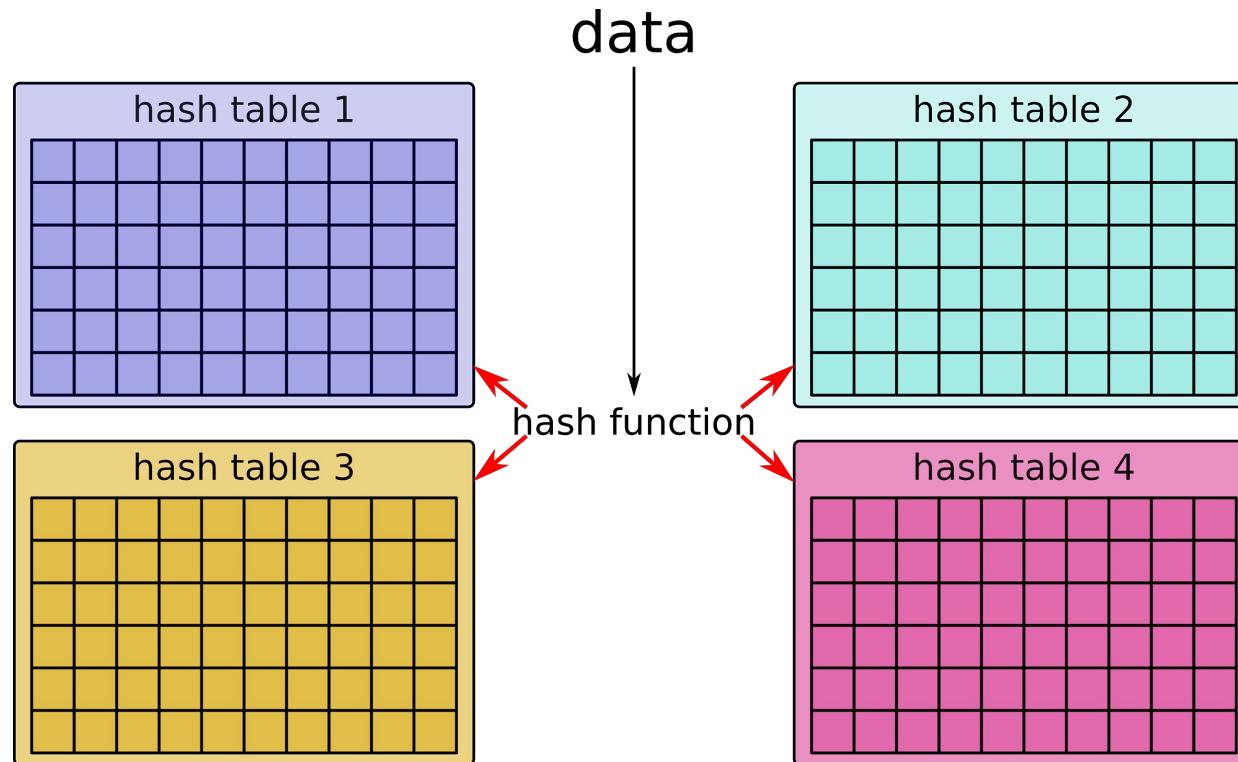
(Don't do this!)



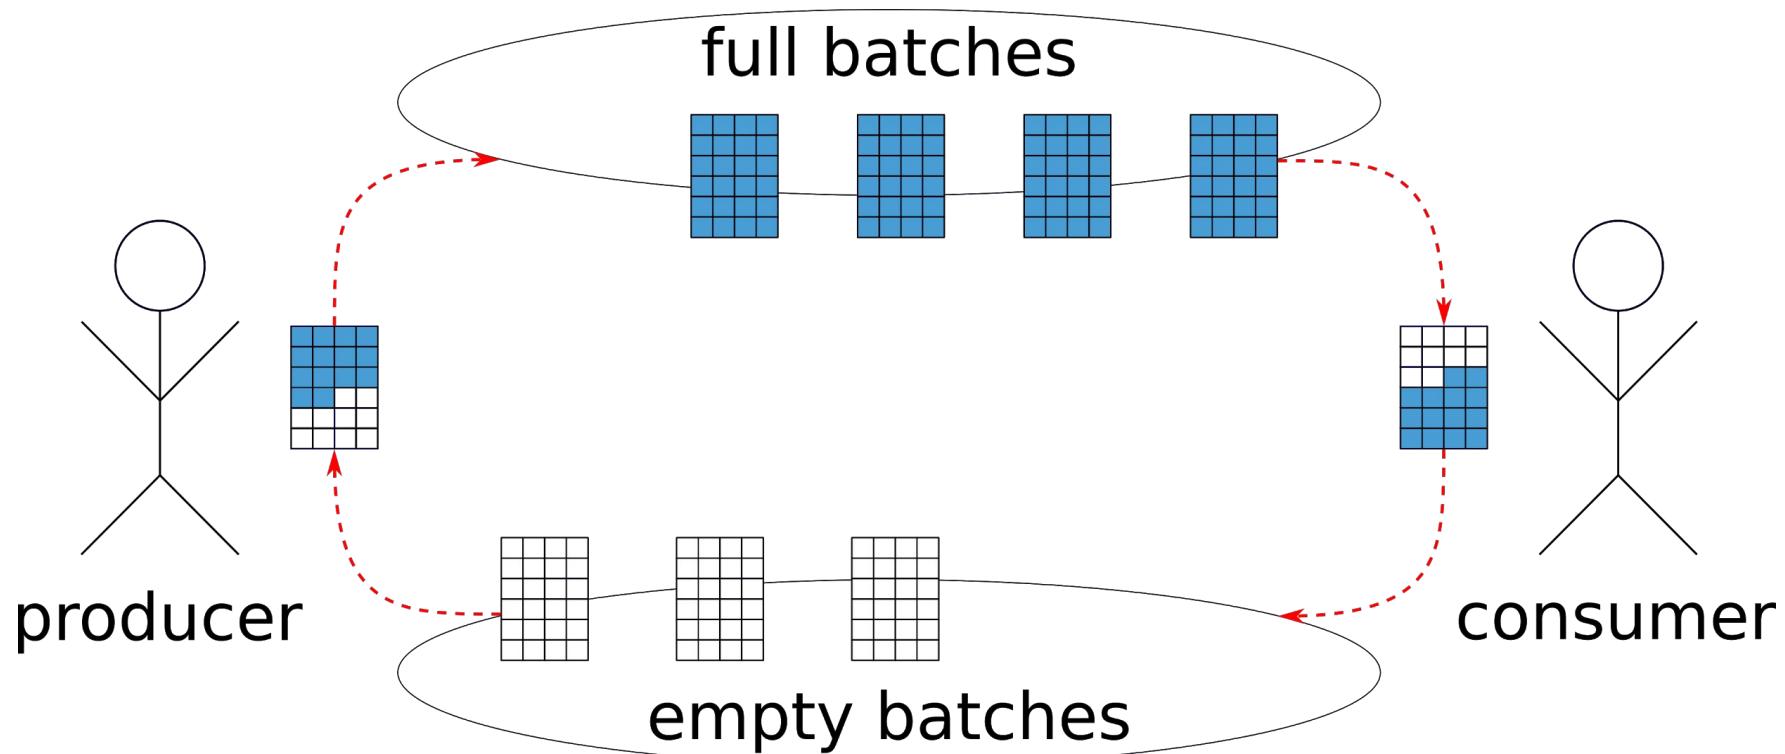
# Atomic compare-and-swap (CAS) instruction

- Can be used to implement **lock-free** algorithms
- **Compare** content of memory location with an expected value
- If the content equals the expected value:
  - Store the new value
  - Return the old value
- One atomic CPU instruction
  - Cannot be interrupted by another thread
- Positions in a hash table are initialized with 0
- Try to store a new key
- Expected value is always 0
- CAS only stores the new key if the slot was empty

# Alternative: Partition hash table into sub-tables



# Producer-consumer model on partitioned table



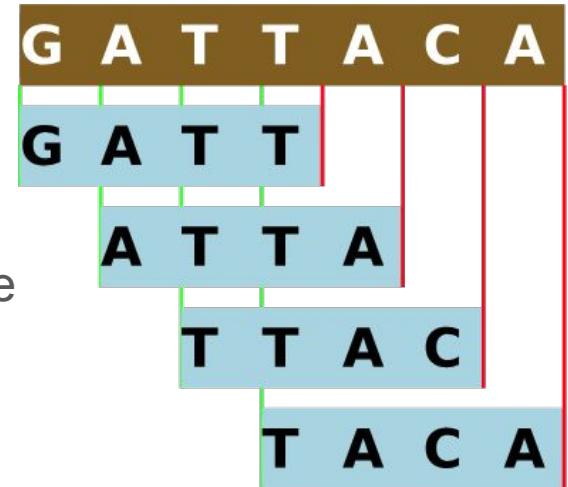
# Summary

# Summary and Recommendations

- Use  $(2, b)$  Cuckoo hashing with  $b = 4$  to 8
- Use  $(3, b)$  Cuckoo hashing with  $b = 2$  to 8
- Ensure each bucket is contained in a cache line (no add'l cache misses!)
- Perhaps use  $(h,b)$ -Cuckoo hashing with overlapping buckets  
(Espinoza: "Cuckoo breeding ground", unpublished)
- All of the above offer high loads and fast access  
(close to 1 lookup per element is possible).

# Xengsort: Alignment-free xenograft sorting

- Partition each read into its  $k$ -mers ( $k \geq 23$ )
- Look up information on each present  $k$ -mer
  - $k$ -mer  $\mapsto$  human | mouse | both
  - $k$ -mer  $\mapsto$  human | human? | mouse | mouse? | both
  - "weak"  $k$ -mers: Hamming-1 neighbor in other genome
- Aggregate  $k$ -mer information into a statement about the entire read, e.g., by majority vote



# Xengsort

## Hashing strategy:

- $(h,b)$ -Cuckoo hashing with  $h=3$  and  $b=4$ 
  - Insertion strategy: Random walk
- Value: 3 bit for host, graft and weak
- achieves up to 99% load

## Page layout:

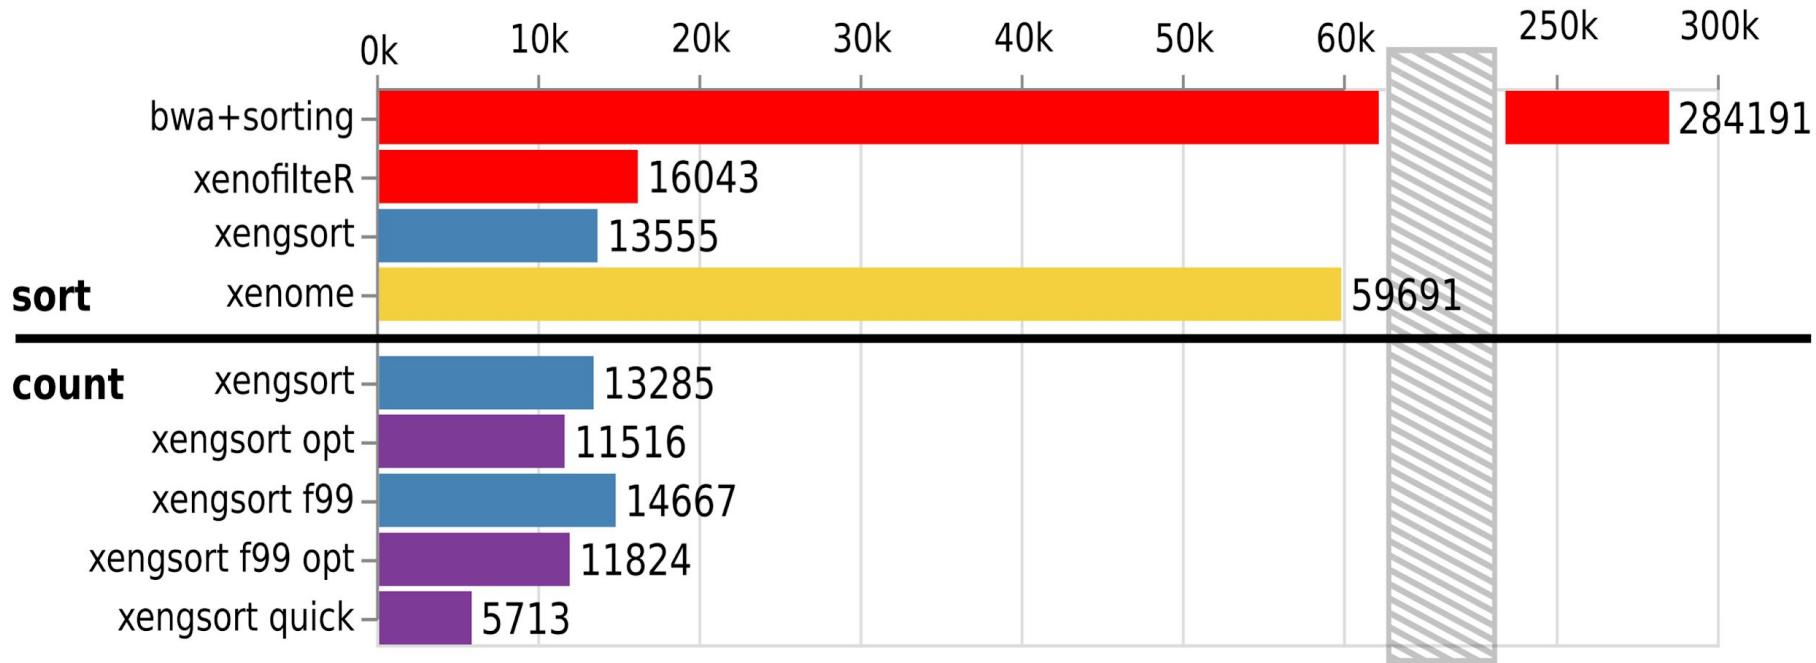
- page = [slot<sub>1</sub>, ..., slot<sub>b</sub>]
- slot = [fingerprint, choice, value]

## Classification:

- Use multiple threads (only read access)
- divide a chunk of reads to each thread
- query  $k$ -mers of each read
- use decision tree to classify reads

# Run time

**Sum of running times [min]**



# Trade-offs

- Larger page sizes (up to one Cacheline)
  - Less pages → **More bits** for the fingerprints
- More hash functions:
  - **more even distribution** of the elements on the pages
  - more positions to check → **more cache misses**